

第六讲

Computer Architecture

ISA Tradeoffs (I)

Prof. Onur Mutlu
Carnegie Mellon University

ISA vs. Microarchitecture Level Tradeoff

- A similar tradeoff (control vs. data-driven execution) can be made at the microarchitecture level
- ISA: Specifies how the programmer sees instructions to be executed
 - Programmer sees a sequential, control-flow execution order vs.
 - Programmer sees a data-flow execution order
- Microarchitecture: How the underlying implementation actually executes instructions
 - Microarchitecture can execute instructions in any order as long as it obeys the semantics specified by the ISA when making the instruction results visible to software
 - Programmer should see the order specified by the ISA

Property of ISA vs. Uarch?

- ADD instruction's opcode
 - Number of general purpose registers
 - Number of ports to the register file
 - Number of cycles to execute the MUL instruction
 - Whether or not the machine employs pipelined instruction execution
-
- Remember
 - Microarchitecture: Implementation of the ISA under specific design constraints and goals

Design Point


- A set of design considerations and their importance
 - ❑ **leads to tradeoffs** in both ISA and uarch
- Considerations
 - ❑ Cost
 - ❑ Performance
 - ❑ Maximum power consumption
 - ❑ Energy consumption (battery life)
 - ❑ Availability
 - ❑ Reliability and Correctness
 - ❑ Time to Market
- Design point determined by the “Problem” space (application space)

Problem
Algorithm
Program
ISA
Microarchitecture
Circuits
Electrons

Tradeoffs: Soul of Computer Architecture

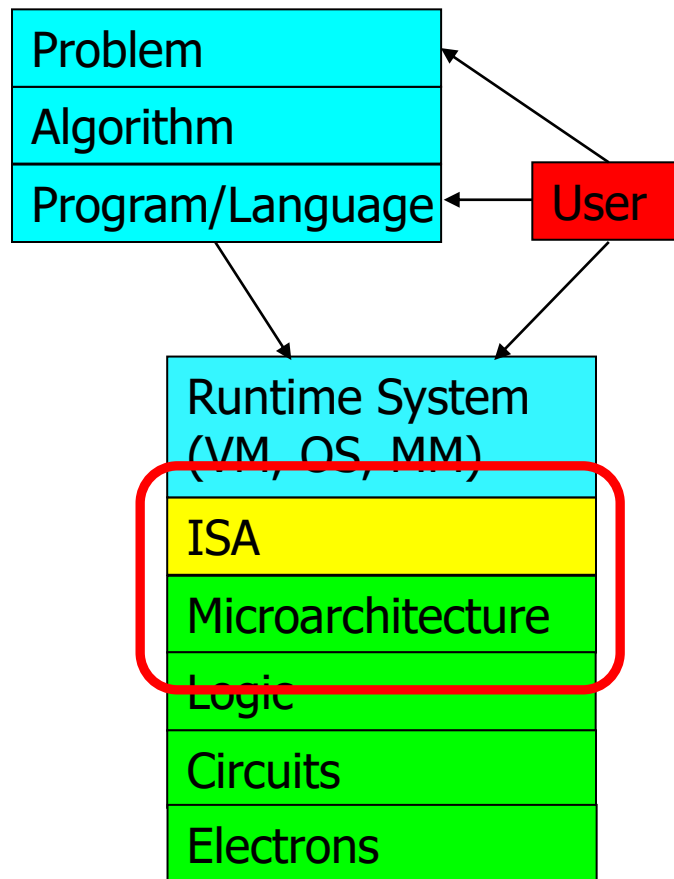
- ISA-level tradeoffs
- Microarchitecture-level tradeoffs
- System and Task-level tradeoffs
 - How to divide the labor between hardware and software
- *Computer architecture is the science and art of making the appropriate trade-offs to meet a design point*
 - *Why art?*

Why Is It (Somewhat) Art?



- We do not (fully) know the future (applications, users, market)

Why Is It (Somewhat) Art?



- And, the future is not constant (it changes)!

ISA Principles and Tradeoffs

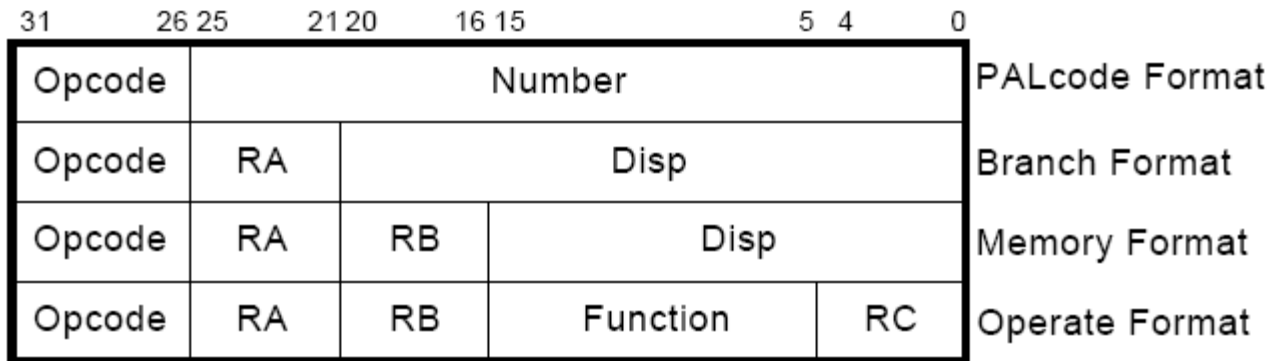
Many Different ISAs Over Decades

- x86
- PDP-x: Programmed Data Processor (PDP-11)
- VAX
- IBM 360
- CDC 6600
- SIMD ISAs: CRAY-1, Connection Machine
- VLIW ISAs: Multiflow, Cydrome, IA-64 (EPIC)
- PowerPC, POWER
- RISC ISAs: Alpha, MIPS, SPARC, ARM

- What are the fundamental differences?
 - E.g., how instructions are specified and what they do
 - E.g., how complex are the instructions

Instruction

- Basic element of the HW/SW interface
- Consists of
 - opcode: what the instruction does
 - operands: who it is to do it to
- Example from Alpha ISA:



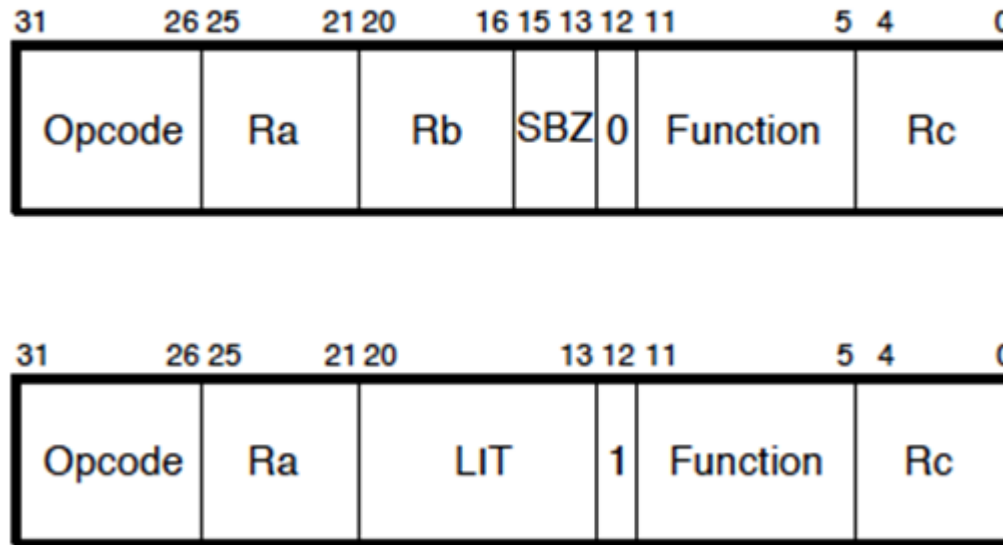
Set of Instructions, Encoding, and Spec

	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
ADD ⁺	0001				DR			SR1			A	op.spec				
AND ⁺	0101				DR			SR1			A	op.spec				
BR	0000				n	z	p	PCOffset9								
JMP	1100				000			BaseR			000000					
JSR(R)	0100				A	operand.specifier										
LDB ⁺	0010				DR			BaseR			boffset6					
LDW ⁺	0110				DR			BaseR			offset6					
LEA ⁺	1110				DR			PCOffset9								
RTI	1000				000000000000											
SHF ⁺	1101				DR			SR			A	D	amount4			
STB	0011				SR			BaseR			boffset6					
STW	0111				SR			BaseR			offset6					
TRAP	1111				0000			trapvect8								
XOR ⁺	1001				DR			SR1			A	op.spec				
not used	1010															
not used	1011															

- Example from LC-3b ISA
 - http://www.ece.utexas.edu/~patt/11s.460N/handouts/new_byte.pdf
- x86 Manual
- Aside: concept of “bit steering”
 - A bit in the instruction determines the interpretation of other bits
- Why unused instructions?

Bit Steering in Alpha

Figure : Operate Instruction Format



If bit <12> of the instruction is 0, the Rb field specifies a source register operand.

If bit <12> of the instruction is 1, an 8-bit zero-extended literal constant is formed by bits <20:13> of the instruction. The literal is interpreted as a positive integer between 0 and 255 and is zero-extended to 64 bits.

What Are the Elements of An ISA?

■ Instruction sequencing model

- Control flow vs. data flow
- Tradeoffs?

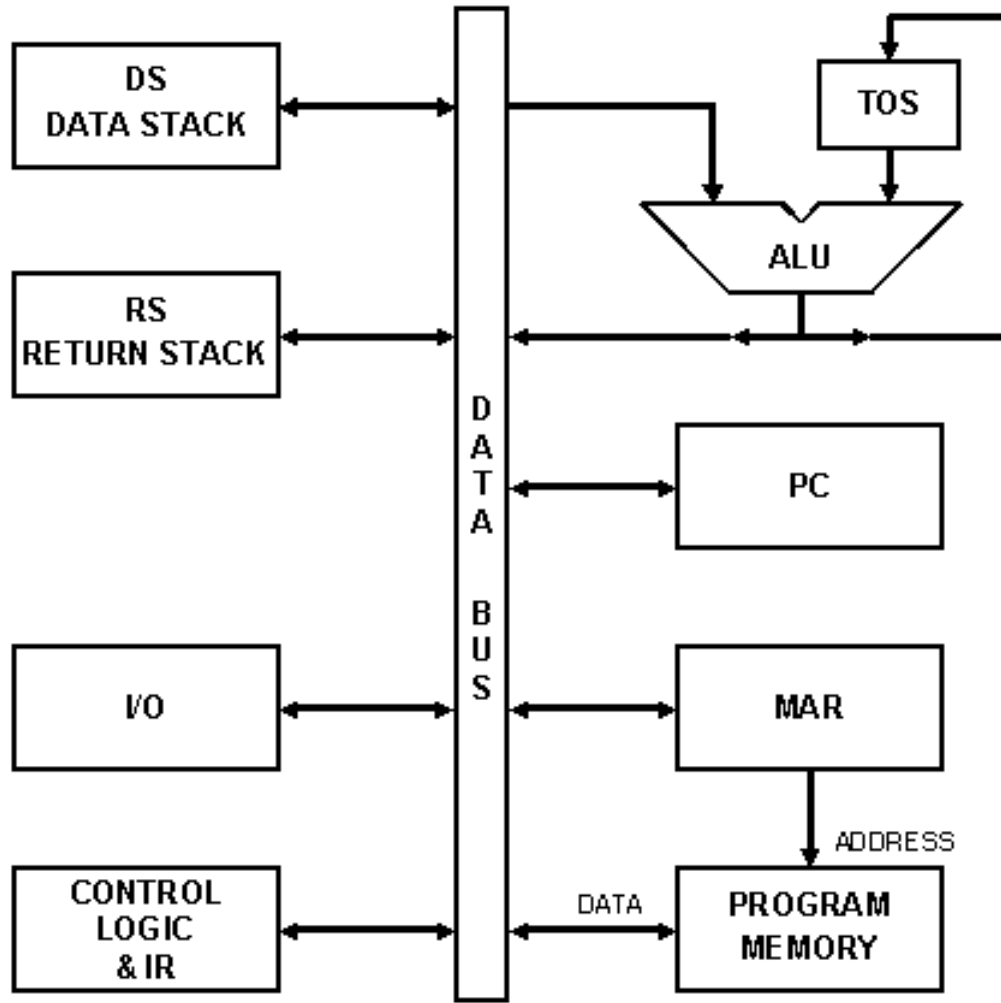
■ Instruction processing style

- Specifies the number of “operands” an instruction “operates” on and how it does so
- 0, 1, 2, 3 address machines
 - 0-address: stack machine (push A, pop A, op)
 - 1-address: accumulator machine (ld A, st A, op A)
 - 2-address: 2-operand machine (one is both source and dest)
 - 3-address: 3-operand machine (source and dest are separate)
- Tradeoffs?
 - Larger operate instructions vs. more executed operations
 - Code size vs. execution time vs. on-chip memory space

An Example: Stack Machine

- + Small instruction size (no operands needed for operate instructions)
 - ❑ Simpler logic
 - ❑ Compact code
- + Efficient procedure calls: all parameters on stack
 - ❑ No additional cycles for parameter passing
- Computations that are not easily expressible with “postfix notation” are difficult to map to stack machines
 - ❑ Cannot perform operations on many values at the same time (only top N values on the stack at the same time)
 - ❑ Not flexible
 - ❑ $(a+b)*c$ 逆波兰表示 $ab+c*$

An Example: Stack Machine (II)




Koopman, “[Stack Computers: The New Wave](http://www.ece.cmu.edu/~koopman/stack_computers/sec3_2.html),” 1989.

http://www.ece.cmu.edu/~koopman/stack_computers/sec3_2.html

Figure 3.1 -- The canonical stack machine.

An Example: Stack Machine Operation



Koopman, “[Stack Computers: The New Wave](http://www.ece.cmu.edu/~koopman/stack_computers/sec3_2.html),” 1989.

http://www.ece.cmu.edu/~koopman/stack_computers/sec3_2.html

运算：
(45+12)*98 后缀表达式 45 12+98*

Figure 3.2 -- An example stack machine.

Other Examples

- PDP-11: A 2-address machine
 - PDP-11 ADD: 4-bit opcode, 2 6-bit operand specifiers
ADD OPRD1,OPRD2
 - Why? Limited bits to specify an instruction
 - Disadvantage: One source operand is always clobbered with the result of the instruction
 - *How do you ensure you preserve the old value of the source?*

- X86: A 2-address (memory/memory) machine
- Alpha: A 3-address (load/store) machine
- MIPS?

What Are the Elements of An ISA?

■ Instructions

- ❑ Opcode
- ❑ Operand specifiers (addressing modes)
 - How to obtain the operand? *Why are there different addressing modes?*

■ Data types

- ❑ Definition: Representation of information for which there are instructions that operate on the representation
- ❑ Integer, floating point, character, binary, decimal, BCD
- ❑ Doubly linked list, queue, string, bit vector, stack
 - VAX: INSQUEUE and REMQUEUE instructions on a doubly linked list or queue; FINDFIRST
 - Digital Equipment Corp., “[VAX11 780 Architecture Handbook](#),” 1977.
 - X86: SCAN opcode operates on character strings; PUSH/POP

Data Type Tradeoffs

- What is the benefit of having more or high-level data types in the ISA?
 - What is the disadvantage?
 - Think compiler/programmer vs. microarchitecture
 - Concept of semantic gap
 - Data types coupled tightly to the semantic level, or complexity of instructions
 - Example: Early RISC architectures vs. Intel 432
 - Early RISC: Only integer data type
 - Intel 432: Object data type, capability based machine
-

What Are the Elements of An ISA?

■ Memory organization

- ❑ Address space: How many uniquely identifiable locations in memory (内存中有多少唯一可识别的位置?)
- ❑ Addressability: How much data does each uniquely identifiable location store (每个唯一可识别的位置能够存储多少数据?)
 - Byte addressable: most ISAs, characters are 8 bits
 - Bit addressable: Burroughs 1700.
 - 64-bit addressable: Some supercomputers.
 - 32-bit addressable: First Alpha
 - Food for thought
 - ❑ How do you add 2 32-bit numbers with only byte addressability?
 - ❑ How do you add 2 8-bit numbers with only 32-bit addressability?
 - ❑ **Big endian vs. little endian?** MSB at low or high byte.
- ❑ Support for virtual memory

Big endian vs. little endian?

- 将一个32位的整数0x12345678存放到一个整型变量(int)(连续地址)中：

地址偏移	Big endian	Little endian
0x00	12 (MSB)	78
0x01	34	56
0x02	56	34
0x03	78	12 (MSB)

Some Historical Readings

- If you want to dig deeper
- Wilner, “Design of the Burroughs 1700,” AFIPS 1972.
- Levy, “The Intel iAPX 432,” 1981.
 - <http://www.cs.washington.edu/homes/levy/capabook/Chapter9.pdf>

What Are the Elements of An ISA?


■ Registers

- How many
- Size of each register


■ Why is having registers a good idea?

- Because programs exhibit a characteristic called data locality
- A recently produced/accessed value is likely to be used more than once (temporal locality)
 - Storing that value in a register eliminates the need to go to memory each time that value is needed

Programmer Visible (Architectural) State




Memory
array of storage locations
indexed by an address



Registers

- given special names in the ISA (as opposed to addresses)
- general vs. special purpose



memory address of the current instruction

Instructions (and programs) specify how to transform
the values of programmer visible state

Aside: Programmer Invisible State

- Microarchitectural state
- Programmer cannot access this directly

- E.g. cache state
- E.g. pipeline registers

Evolution of Register Architecture

- Accumulator
 - a legacy from the “adding” machine days
- Accumulator + address registers
 - need register indirection
 - initially address registers were special-purpose, i.e., can only be loaded with an address for indirection
 - eventually arithmetic on addresses became supported
- General purpose registers (GPR)
 - all registers good for all purposes
 - grew from a few registers to 32 (common for RISC) to 128 in Intel IA-64

Instruction Classes

- Operate instructions
 - Process data: arithmetic and logical operations
 - Fetch operands, compute result, store result
 - Implicit sequential control flow
- Data movement instructions
 - Move data between memory, registers, I/O devices
 - Implicit sequential control flow
- Control flow instructions
 - Change the sequence of instructions that are executed

What Are the Elements of An ISA?

- Load/store vs. memory/memory architectures
 - Load/store architecture: operate instructions operate only on registers
 - E.g., MIPS, ARM and many RISC ISAs
 - Memory/memory architecture: operate instructions can operate on memory locations
 - E.g., x86, VAX and many CISC ISAs

What Are the Elements of An ISA?

■ Addressing modes specify how to obtain the operands

- ❑ Absolute LW rt, 10000
use immediate value as address
- ❑ Register Indirect: LW rt, (r_{base})
use $\text{GPR}[r_{\text{base}}]$ as address
- ❑ Displaced or based: LW rt, offset(r_{base})
use $\text{offset} + \text{GPR}[r_{\text{base}}]$ as address
- ❑ Indexed: LW rt, ($r_{\text{base}}, r_{\text{index}}$)
use $\text{GPR}[r_{\text{base}}] + \text{GPR}[r_{\text{index}}]$ as address
- ❑ Memory Indirect LW rt ((r_{base}))
use value at $\text{M}[\text{GPR}[r_{\text{base}}]]$ as address
- ❑ Auto inc/decrement LW Rt, (r_{base})
use $\text{GRP}[r_{\text{base}}]$ as address, but inc. or dec. $\text{GPR}[r_{\text{base}}]$ each time

What Are the Benefits of Different Addressing Modes?

- Another example of programmer vs. microarchitect tradeoff
- Advantage of more addressing modes:
 - Enables better mapping of high-level constructs to the machine: some accesses are better expressed with a different mode → reduced number of instructions and code size
 - Think array accesses (autoincrement mode)
 - Think indirection (pointer chasing)
 - Sparse matrix accesses
- Disadvantage:
 - More work for the compiler
 - More work for the microarchitect

ISA Orthogonality

- Orthogonal ISA: (正交ISA: 寻址模式x操作x数据类型)
 - All addressing modes can be used with all instruction types
 - Example: VAX
 - (~13 addressing modes) x (>300 opcodes) x (integer and FP formats)
- Who is this good for?
- Who is this bad for?

What Are the Elements of An ISA?

- How to interface with I/O devices
 - Memory mapped I/O
 - A region of memory is mapped to I/O devices
 - I/O operations are loads and stores to those locations
 - Special I/O instructions
 - IN and OUT instructions in x86 deal with ports of the chip
 - Tradeoffs?
 - Which one is more general purpose?

What Are the Elements of An ISA?

■ Privilege modes

- User vs supervisor (MIPS:核心态, 超级用户态, 用户态)
- Who can execute what instructions?

■ Exception and interrupt handling

- What procedure is followed when something goes wrong with an instruction?
- What procedure is followed when an external device requests the processor?
- Vectored vs. non-vectored interrupts (early MIPS)

■ Virtual memory

- Each program has the illusion of the entire memory space, which is greater than physical memory

■ Access protection

MIPS exceptions 举例

Table 优先级异常入口

异常类型	正常运行 (BEV 为 0)	启动 (BEV 为 1)
冷启动、热重启、非屏蔽中断	0xFFFFFFFF BFC00000	0xFFFFFFFF BFC00000
TLB 重填	0xFFFFFFFF 80000000	0xFFFFFFFF BFC00200
xTLB 重填	0xFFFFFFFF 80000080	0xFFFFFFFF BFC00280
cache 错误	0xFFFFFFFF A0000100	0xFFFFFFFF BFC00300
其他	0xFFFFFFFF 80000180	0xFFFFFFFF BFC00380