





Applications of Parallel Computers

Data Parallel Algorithms

<https://sites.google.com/lbl.gov/cs267-spr2021>



Parallel Machines and Programming



Abstract
Machine
Models

Shared Memory

Processors execute own instruction stream
Communicate by reading/writing memory

Distributed Memory

Processors execute own instruction stream
Communicate by sending messages

Single Instruction Multiple Data (SIMD)

One instruction stream (all run same instruction)
Communicate through memory

The Power of Data Parallelism

- Data parallelism: perform the same operation on multiple values (often array elements)
 - Also includes reductions, broadcast, scan..
- Many parallel programming models use some data parallelism
 - SIMD units (and previously SIMD supercomputers)
 - CUDA / GPUs
 - MapReduce
 - MPI collectives

Data Parallel Programming: Unary Operators

- Unary operations applied to all elements of an array

A = array

B = array

f = square (any unary function, i.e., 1 argument)

B = f(A)

A:

3	1	1	2	3	3	4	2	2	2	1	3	1	1	1	3	3	2	1
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---



f applied to each element

B:

9	1	1	4	9	9	16	4	4	4	1	9	1	1	1	9	9	4	1
---	---	---	---	---	---	----	---	---	---	---	---	---	---	---	---	---	---	---

Data Parallel Programming: Binary Operators

- Binary operations applied to all pairs of elements

A = array

B = array

C = array

- or any other binary operator

C = A - B

A:

3	1	0	2	3	0	4	2	0	2	1	3	0	1	1	0	3	2	1
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

- applied to each pair

B:

0	1	1	4	1	0	2	1	4	3	1	0	1	1	2	3	5	3	2
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---




C:

3	0	-1	-2	2	0	2	2	-4	-2	0	3	-1	0	-1	-3	-2	-1	-1
---	---	----	----	---	---	---	---	----	----	---	---	----	---	----	----	----	----	----

Data Parallel Programming: Broadcast

- Broadcast fill a value into all elements of an array



- Useful for $a*X+Y$ called axpy, saxpy, daxpy
 - For single, double precision, or in general

Memory Operations: Strided and Scatter / Gather

- Array assignment works if the arrays are the same shape

A: double [0:4]

B: double [0:4] = [0.0, 1.0, 2.2, 3.3, 4.4]

A = B

Memory Operations: Strided and Scatter / Gather

- Array assignment works if the arrays are the same shape

A: double [0:4]

B: double [0:4] = [0.0, 1.0, 2.2, 3.3, 4.4]

A = B

- May have a stride, i.e., not be contiguous in memory

A = B [0:4:2] // copy with stride 2 (every other element)

C: double [0:4, 0:4]

A = C[:,3] // copy column of C

Memory Operations: Strided and Scatter / Gather

- Array assignment works if the arrays are the same shape

A: double [0:4]

B: double [0:4] = [0.0, 1.0, 2.2, 3.3, 4.4]

A = B

- May have a stride, i.e., not be contiguous in memory

A = B [0:4:2] // copy with stride 2 (every other element)

C: double [0:4, 0:4]

A = C[:,3] // copy column of C

- Gather (indexed) values from one array

X: int [0:4] = [3, 0, 4, 2, 1] // a permutation of indices 0 to 4

A = B[X] // A now is [3.3, 0.0, 4.4, 2.2, 1.1]

Memory Operations: Strided and Scatter / Gather

- Array assignment works if the arrays are the same shape

A: double [0:4]

B: double [0:4] = [0.0, 1.0, 2.2, 3.3, 4.4]

A = B

- May have a stride, i.e., not be contiguous in memory

A = B [0:4:2] // copy with stride 2 (every other element)

C: double [0:4, 0:4]

A = C[:,3] // copy column of C

- Gather (indexed) values from one array

X: int [0:4] = [3, 0, 4, 2, 1] // a permutation of indices 1- 4

A = B[X] // A now is [3.3, 0.0, 4.4, 2.2, 1.1]

- Scatter (indexed) values from one array

A[X] = B // A now is [1.1, 4.4, 3.3, 0.0, 2.2]

Memory Operations: Strided and Scatter / Gather

- Array assignment works if the arrays are the same shape

A: double [0:4]

B: double [0:4] = [0.0, 1.0, 2.2, 3.3, 4.4]

A = B

- May have a stride, i.e., not be contiguous in memory

A = B [0:4:2] // copy with stride 2 (every other element)

C: double [0:4, 0:4]

A = C[:,3] // copy column of C

- Gather (indexed) values from one array

X: int [0:4] = [3, 0, 4, 2, 1] // a permutation of indices 1- 4

A = B[X] // A now is [3.3, 0.0, 4.4, 2.2, 1.1]

- Scatter (indexed) values from one array

A[X] = B // A now is [1.1, 4.4, 3.3, 0.0, 2.2]

Data Parallel Programming: Masks

- Can apply operations under a “mask”

M = array of 0/1 (True/False)

A = array

B = array

A = A + B under M

A:

3	1	1	2	3	3	4	2	2	2	1	3	1	1	1	3	3	2	1
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

B:

0	1	1	4	1	0	2	1	4	3	1	0	1	1	2	3	5	3	2
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

M:

1	0	0	1	1	0	0	0	1	1	1	1	0	0	1	0	0	1	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---




+ under mask

A:

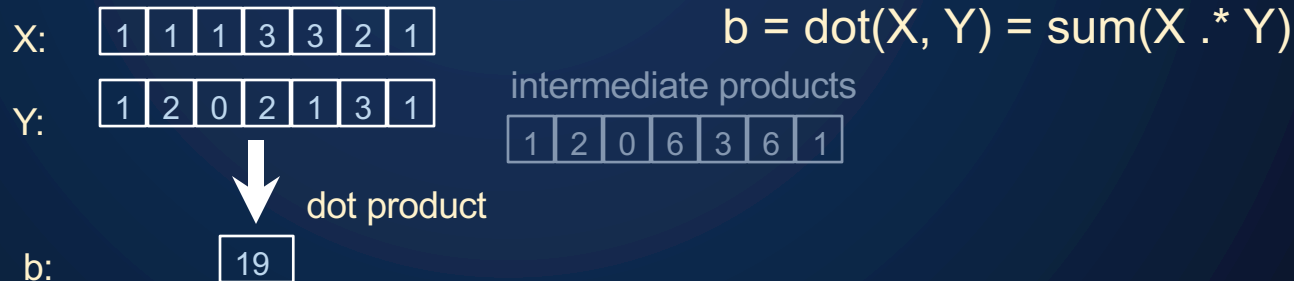
0			4	1				4	3	1	0			2			3	
3	1	1	6	4	3	4	2	6	5	2	3	1	1	3	3	3	5	1

Data Parallel Programming: Reduce

- Reduce an array to a value with + or any associative op




- Associative so we can perform op in different order
- Useful for dot products (ddot, sdot, etc.) $b = X^T Y = \sum_j X[j] * Y[j]$



Data Parallel Programming: Scans


- Fill array with partial reductions any associative op
- Sum scan:

A = array
B = array
B = **scan(A,+)**



- Max scan:

B = **scan(A,max)**



Inclusive and Exclusive Scans

Two variations of a scan, given an input vector $[x_0, x_1, \dots, x_{n-1}]$:

- *inclusive* scan includes input x_i when computing output y_i

$$[a_0, (a_0 \odot a_1), \dots, (a_0 \odot a_1 \dots \odot a_{n-1})]$$

e.g., `add_scan_inclusive([1, 0, 3, 0, 2])` \rightarrow `[1, 1, 4, 4, 6]`

- *exclusive* scan does *not* x_i when computing output y_i

$$[I, a_0, (a_0 \odot a_1), \dots, (a_0 \odot a_1 \dots \odot a_{n-2})] \text{ where } I \text{ is the identity for } \odot$$

e.g., `add_scan_exclusive([1, 0, 3, 0, 2])` \rightarrow `[0, 1, 1, 4, 4]`

Inclusive and Exclusive Scans

Two variations of a scan, given an input vector $[x_0, x_1, \dots, x_{n-1}]$:

- *inclusive* scan includes input x_i when computing output y_i

$$[a_0, (a_0 \odot a_1), \dots, (a_0 \odot a_1 \dots \odot a_{n-1})]$$

e.g., `add_scan_inclusive([1, 0, 3, 0, 2])` \rightarrow `[1, 1, 4, 4, 6]`

- *exclusive* scan does *not* x_i when computing output y_i

$$[I, a_0, (a_0 \odot a_1), \dots, (a_0 \odot a_1 \dots \odot a_{n-2})] \text{ where } I \text{ is the identity for } \odot$$

e.g., `add_scan_exclusive([1, 0, 3, 0, 2])` \rightarrow `[0, 1, 1, 4, 4]`

Can easily get the inclusive version from the exclusive:

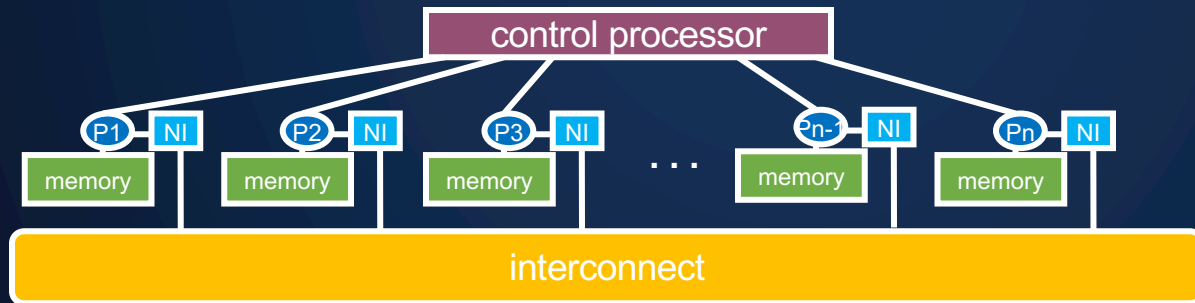
$$\text{scan_inclusive}(X) = X \odot \text{scan_exclusive}(X).$$

For the other way you need an inverse for \odot (or shift)

Idealized Hardware and Performance Model

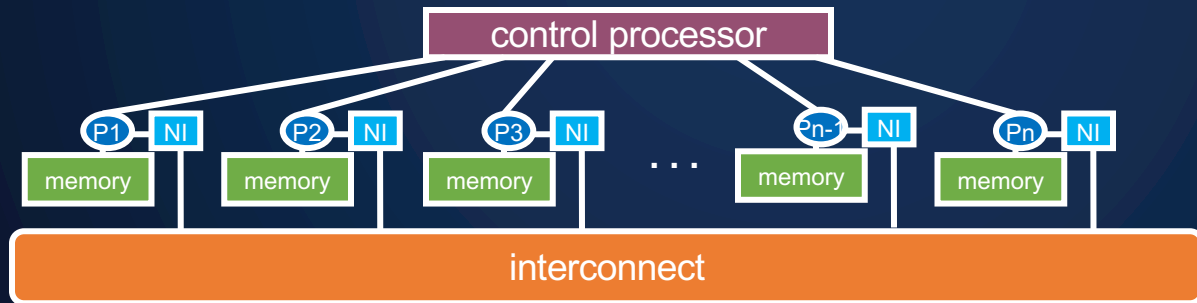
SIMD Systems Implemented Data Parallelism

- SIMD Machine: A large number of (usually) tiny processors.
 - A single “control processor” issues each instruction.
 - Each processor executes the same instruction.
 - Some processors may be turned off on some instructions.




Ideal Cost Model for Data Parallelism

- Machine
 - An unbounded number of processors (p)
 - Control overhead is free
 - Communication is free
- Cost (complexity) on this abstract machine is the algorithm's *span or depth*, T_∞
 - Defines a lower bound on time on real machines




Cost on Ideal Machine (Span)

- Span for unary or binary operations (pleasingly parallel)



- Even if arrays are not aligned, communication is “free” here
- Reductions and broadcasts



Broadcast and reduction on processor tree

- Broadcast of 1 value to p processors with $\log n$ span



Broadcast



Add-reduction

- Reduction of n values to 1 with $\log n$ span
- Takes advantage of associativity in $+$, $*$, \min , \max , etc.

Can reductions go faster? No, $\log n$ lower bound
on any function of n variables!

n “useful” inputs




Can reductions go faster? No, $\log n$ lower bound on any function of n variables!

- Given a function $f(x_1, \dots, x_n)$ of n input variables and 1 output variable, how fast can we evaluate it in parallel?
- Assume we only have binary operations, one per time step
- After 1 time step, an output can only depend on two inputs




Can reductions go faster? No, $\log n$ lower bound on any function of n variables!

- Given a function $f(x_1, \dots, x_n)$ of n input variables and 1 output variable, how fast can we evaluate it in parallel?
- Assume we only have binary operations, one per time step
- After 1 time step, an output can only depend on two inputs
- By induction: after k time units, an output can only depend on 2^k inputs
 - After $\log_2 n$ time units, output depends on at most n inputs
- A binary tree performs such a computation



Multiplying n-by-n matrices in $O(\log n)$ time

- Use n^3 processors
- Step 1: For all $(1 \leq i, j, k \leq n)$ $P(i, j, k) = A(i, k) * B(k, j)$
 - cost = 1 time unit, using n^3 processors
- Step 2: For all $(1 \leq i, j \leq n)$ $C(i, j) = \sum_{k=1}^n P(i, j, k)$
 - cost = $O(\log n)$ time, using n^2 trees, $n^3 / 2$ processors each



Put a
processor at
every point in
this cube

What about Scan (aka Parallel Prefix)?

- Recall: the **scan** operation takes a **binary associative** operator \odot , and an array of n elements

$[a_0, a_1, a_2, \dots, a_{n-1}]$

and produces the array

$[a_0, (a_0 \odot a_1), \dots, (a_0 \odot a_1 \odot \dots \odot a_{n-1})]$

- Example: **add scan** of

$[1, 2, 0, 4, 2, 1, 1, 3]$ is $[1, 3, 3, 7, 9, 10, 11, 14]$

- Other operators
 - Reals: $+$, $*$, \min , \max (in floating point will assume associative)
 - Booleans: and , or
 - Matrices: mat mul

Can we parallelize a scan?

- It looks like this:

$y(0) = 0;$

for $i = 1:n$

$y(i) = y(i-1) + x(i);$

- Takes $n-1$ operations (adds) to do in serial

Can we parallelize a scan?

- It looks like this:

```
y(0) = 0;  
for i = 1:n  
    y(i) = y(i-1) + x(i);
```

- Takes $n-1$ operations (adds) to do in serial
- The i^{th} iteration of the loop depends completely on the $(i-1)^{\text{st}}$ iteration.
- Impossible to parallelize, right?

A clue

input = (1, 2, 3, 4, 5, 6, 7, 8)

output = (1, 3, 6, 10, 15, 21, 28, 36)

What if we add, say, $5+6+7+8$?

Parallel But Terribly Inefficient

input = (1, 2, 3, 4, 5, 6, 7, 8)

output = (1, 3, 6, 10, 15, 21, 28, 36)

Put 1 processor at element 1, 2 at element 2, 3 at position 3 ...

- $O(\log n)$ span ☺
- $O(n^2)$ work ☹

A clue

input = (1, 2, 3, 4, 5, 6, 7, 8)

output = (1, 3, 6, 10, 15, 21, 28, 36)

Is there any value in adding, say, $5+6+7+8$?

If we separately have $1+2+3+4$, what can we do?

A clue

input = (1, 2, 3, 4, 5, 6, 7, 8)

output = (1, 3, 6, 10, 15, 21, 28, 36)

Is there any value in adding, say, $5+6+7+8$?

If we separately have $1+2+3+4$, what can we do?

Suppose we added $1+2$, $3+4$, etc. pairwise, is this useful?

Sum Scan (aka prefix sum) in parallel

Algorithm: 1. Pairwise sum 2. Recursive prefix 3. Pairwise sum



Sum Scan (aka prefix sum) in parallel

Algorithm: 1. Pairwise sum 2. Recursive prefix 3. Pairwise sum

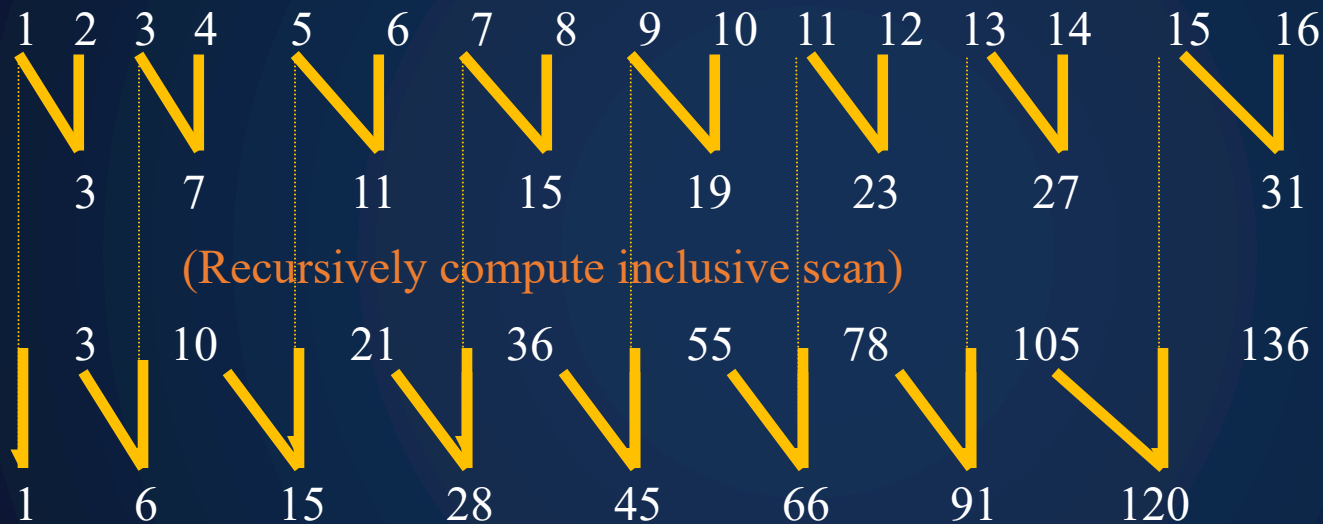


(Recursively compute inclusive scan)

Input 1	Input 2	Pairwise Sum
3	10	21
36	55	78
105	136	


Sum Scan (aka prefix sum) in parallel

Algorithm: 1. Pairwise sum 2. Recursive prefix 3. Pairwise sum



Sum Scan (aka prefix sum) in parallel

Algorithm: 1. Pairwise sum 2. Recursive prefix 3. Pairwise sum




Parallel prefix cost

Time for this algorithm on one processor (work)

- $T_1(n) = n/2 + n/2 + T_1(n/2) = n + T_1(n/2) = 2n - 1$

Time on unbounded number of processors (span)

- $T_\infty(n) = 2 \log n$



Pairwise sum

Recursive prefix

Pairwise sum
(update odds)

Parallelism at the cost of more work (2x)!


Non-recursive exclusive scan

Up-sweep


$d=0$




Non-recursive exclusive scan




Non-recursive exclusive scan




Non-recursive exclusive scan



Non-recursive exclusive scan




Non-recursive exclusive scan



Non-recursive exclusive scan




Non-recursive exclusive scan



Algorithm due to Blelloch

Non-recursive exclusive scan



*This is both
work-efficient
(n adds) and
space-efficient
(update in
place)*

(Non-trivial) Applications of Data Parallelism

..using scans

Scans are useful for many things (partial list here)

- Reduction and broadcast in $O(\log n)$ time
- Parallel prefix (scan) in $O(\log n)$ time
- Adding two n -bit integers in $O(\log n)$ time
- Multiplying n -by- n matrices in $O(\log n)$ time
- Inverting n -by- n triangular matrices in $O(\log^2 n)$ time
- Inverting n -by- n dense matrices in $O(\log^2 n)$ time
- Evaluating arbitrary expressions in $O(\log n)$ time
- Evaluating recurrences in $O(\log n)$ time
- “2D parallel prefix”, for image segmentation (Catanzaro & Keutzer)
- Sparse-Matrix-Vector-Multiply (SpMV) using Segmented Scan
- Parallel page layout in a browser (Leo Meyerovich, Ras Bodik)
- Solving n -by- n tridiagonal matrices in $O(\log n)$ time
- Traversing linked lists
- Computing minimal spanning trees
- Computing convex hulls of point sets...

Scans are useful for many things (partial list here)

- Reduction and broadcast in $O(\log n)$ time
- Parallel prefix (scan) in $O(\log n)$ time
- Adding two n -bit integers in $O(\log n)$ time
- Multiplying n -by- n matrices in $O(\log n)$ time
- Inverting n -by- n triangular matrices in $O(\log^2 n)$ time
- Inverting n -by- n dense matrices in $O(\log^2 n)$ time
- Evaluating arbitrary expressions in $O(\log n)$ time
- Evaluating recurrences in $O(\log n)$ time
- “2D parallel prefix”, for image segmentation (Catanzaro & Keutzer)
- Sparse-Matrix-Vector-Multiply (SpMV) using Segmented Scan
- Parallel page layout in a browser (Leo Meyerovich, Ras Bodik)
- Solving n -by- n tridiagonal matrices in $O(\log n)$ time
- Traversing linked lists
- Computing minimal spanning trees
- Computing convex hulls of point sets...

Application: Stream Compression

- Given an array of 0/1 flags

flags =

1	0	1	1	0	0	1	1
---	---	---	---	---	---	---	---

and an array (stream) of values

values =

3	2	4	1	5	3	3	1
---	---	---	---	---	---	---	---

compress into

result =

3	4	1	3	1
---	---	---	---	---

Application: Stream Compression

- Given an array of 0/1 flags

flags =

1	0	1	1	0	0	1	1
---	---	---	---	---	---	---	---

and an array (stream) of values

values =

3	2	4	1	5	3	3	1
---	---	---	---	---	---	---	---

compress into

result =

3	4	1	3	1
---	---	---	---	---

- Step 1: Compute an exclusive add scan of flags:

index =

0	1	1	2	3	3	3	4
---	---	---	---	---	---	---	---

Application: Stream Compression

- Given an array of 0/1 flags

flags =

1	0	1	1	0	0	1	1
---	---	---	---	---	---	---	---

and an array (stream) of values

values =

3	2	4	1	5	3	3	1
---	---	---	---	---	---	---	---

compress into

result =

3	4	1	3	1
---	---	---	---	---

- Step 1: Compute an exclusive add scan of flags:

index =

0	1	1	2	3	3	3	4
---	---	---	---	---	---	---	---

- Step 2: “Scatter” values into result at index, masked by flags

result[index] = values at flags



Remove matching elements

- Given an array of values, and an int x, remove all elements that are not divisible by x:

```
int find (int x, int y) (y % x == 0) ? 1 : 0;
```

values =

3	5	6	12	4	2	3	0
---	---	---	----	---	---	---	---

flags = apply(values, find)

1	0	1	1	0	0	1	1
---	---	---	---	---	---	---	---

Use previous solution to remove those not divisible

Application: Radix Sort (serial algorithm)

4	7	2	6	3	5	1	0
---	---	---	---	---	---	---	---

Idea: Sort 1 bit at a time:

- 0s on left, 1s on right

Use a “stable” sort:

- Keep order as it, unless things *need to* switch based on the current bit

Start with least-significant bit

- And move up

Application: Radix Sort (serial algorithm)

4	7	2	6	3	5	1	0
---	---	---	---	---	---	---	---

Sort on least significant bit (Bit_0 in $[\text{Bit}_2, \text{Bit}_1, \text{Bit}_0]$)

$XX0 < XX1$ (evens before odds)

Application: Radix Sort (serial algorithm)

4	7	2	6	3	5	1	0
---	---	---	---	---	---	---	---

4	2	6	0	7	3	5	1
---	---	---	---	---	---	---	---

Bit₀=0

Bit₀=1

Sort on least significant bit (Bit₀ in [Bit₂, Bit₁, Bit₀])

XX0 < XX1 (evens before odds)

Application: Radix Sort (serial algorithm)

4	7	2	6	3	5	1	0
---	---	---	---	---	---	---	---

4	2	6	0	7	3	5	1
---	---	---	---	---	---	---	---

Bit₀=0

Bit₀=1

Sort on least significant bit (Bit₀ in [Bit₂, Bit₁, Bit₀])

XX0 < XX1 (evens before odds)

4	2	6	0	7	3	5	1
---	---	---	---	---	---	---	---

4	0	5	1	2	6	7	3
---	---	---	---	---	---	---	---

Bit₁=0

Bit₁=1

Stably sort entire array on next bit

X0X < X1X

Application: Radix Sort (serial algorithm)

4	7	2	6	3	5	1	0
---	---	---	---	---	---	---	---

Sort on least significant bit (Bit_0 in $[\text{Bit}_2, \text{Bit}_1, \text{Bit}_0]$)

4	2	6	0	7	3	5	1
---	---	---	---	---	---	---	---

$\text{Bit}_0=0$

$\text{Bit}_0=1$

$XX0 < XX1$ (evens before odds)

4	2	6	0	7	3	5	1
---	---	---	---	---	---	---	---

Stably sort entire array on next bit

$X0X < X1X$

Application: Radix Sort (serial algorithm)

4	7	2	6	3	5	1	0
---	---	---	---	---	---	---	---

Sort on least significant bit (Bit_0 in $[\text{Bit}_2, \text{Bit}_1, \text{Bit}_0]$)

4	2	6	0	7	3	5	1
---	---	---	---	---	---	---	---

$\text{Bit}_0=0$

$\text{Bit}_0=1$

$XX0 < XX1$ (evens before odds)

4	2	6	0	7	3	5	1
---	---	---	---	---	---	---	---

Stably sort entire array on next bit

$X0X < X1X$

Application: Radix Sort (serial algorithm)

4	7	2	6	3	5	1	0
---	---	---	---	---	---	---	---

Sort on least significant bit (Bit_0 in $[\text{Bit}_2, \text{Bit}_1, \text{Bit}_0]$)

4	2	6	0	7	3	5	1
---	---	---	---	---	---	---	---

$\text{Bit}_0=0$

$\text{Bit}_0=1$

$XX0 < XX1$ (evens before odds)

4	2	6	0	7	3	5	1
---	---	---	---	---	---	---	---

Stably sort entire array on next bit

4	0	5	1	2	6	7	3
---	---	---	---	---	---	---	---

$\text{Bit}_1=0$

$\text{Bit}_1=1$

$X0X < X1X$

Application: Radix Sort (serial algorithm)

4	7	2	6	3	5	1	0
---	---	---	---	---	---	---	---

Sort on least significant bit (Bit_0 in $[\text{Bit}_2, \text{Bit}_1, \text{Bit}_0]$)

4	2	6	0	7	3	5	1
---	---	---	---	---	---	---	---

$\text{Bit}_0=0$

$\text{Bit}_0=1$

$XX0 < XX1$ (evens before odds)

4	2	6	0	7	3	5	1
---	---	---	---	---	---	---	---

Stably sort entire array on next bit

4	0	5	1	2	6	7	3
---	---	---	---	---	---	---	---

$\text{Bit}_1=0$

$\text{Bit}_1=1$

$X0X < X1X$

4	0	5	1	2	6	7	3
---	---	---	---	---	---	---	---

Stably sort on next bit

$0XX < 1XX$ (<4 before ≥ 4 here)

Application: Radix Sort (serial algorithm)

4	7	2	6	3	5	1	0
---	---	---	---	---	---	---	---

Sort on least significant bit (Bit_0 in $[\text{Bit}_2, \text{Bit}_1, \text{Bit}_0]$)

4	2	6	0	7	3	5	1
---	---	---	---	---	---	---	---

$\text{Bit}_0=0$

$\text{Bit}_0=1$

$XX0 < XX1$ (evens before odds)

4	2	6	0	7	3	5	1
---	---	---	---	---	---	---	---

Stably sort entire array on next bit

4	0	5	1	2	6	7	3
---	---	---	---	---	---	---	---

$X0X < X1X$

$\text{Bit}_1=0$

$\text{Bit}_1=1$

4	0	5	1	2	6	7	3
---	---	---	---	---	---	---	---

Stably sort on next bit

$0XX < 1XX$ (<4 before ≥ 4 here)

Application: Radix Sort (serial algorithm)

4	7	2	6	3	5	1	0
---	---	---	---	---	---	---	---

4	2	6	0	7	3	5	1
---	---	---	---	---	---	---	---

Bit₀=0

Bit₀=1

Sort on least significant bit (Bit₀ in [Bit₂, Bit₁, Bit₀])

XX0 < XX1 (evens before odds)

4	2	6	0	7	3	5	1
---	---	---	---	---	---	---	---

4	0	5	1	2	6	7	3
---	---	---	---	---	---	---	---

Bit₁=0

Bit₁=1

Stably sort entire array on next bit

X0X < X1X

4	0	5	1	2	6	7	3
---	---	---	---	---	---	---	---

0	1	2	3	4	5	6	7
---	---	---	---	---	---	---	---

Bit₂=0

Bit₂=1

Stably sort on next bit

0XX < 1XX (<4 before >=4 here)

Application: Data Parallel Radix Sort

4	7	2	6	3	5	1	0
---	---	---	---	---	---	---	---

input

odds = last bit of each element

evens = complement of odds (last bit = 0)

evpos = exclusive sum scans of evens

totalEvens = broadcast last element

indx = constant array of 0..n

oddpos = totalEvens + indx - evpos

pos = if evens then evpos else oddpos

Scatter input using pos as index

Repeat with next bit to left until done

This will just
do one step
of radix sort
(a stable sort
on 1 bit)

Application: Data Parallel Radix Sort

4	7	2	6	3	5	1	0
0	1	0	0	1	1	1	0

input

odds = last bit of each element

evens = complement of odds (last bit = 0)

evpos = exclusive sum scans of evens

totalEvens = broadcast last element

indx = constant array of 0..n

oddpos = totalEvens + indx - epos

pos = if evens then evpos else oddpos

Scatter input using pos as index

Repeat with next bit to left until done

Application: Data Parallel Radix Sort

4	7	2	6	3	5	1	0
0	1	0	0	1	1	1	0

input

odds = last bit of each element

evens = complement of odds (last bit = 0)

evpos = exclusive sum scans of evens

totalEvens = broadcast last element

indx = constant array of 0..n

oddpos = totalEvens + indx - evpos

pos = if evens then evpos else oddpos

Scatter input using pos as index

Repeat with next bit to left until done

Application: Data Parallel Radix Sort

4	7	2	6	3	5	1	0
0	1	0	0	1	1	1	0
1	0	1	1	0	0	0	1

input

odds = last bit of each element

evens = complement of odds (last bit = 0)

evpos = exclusive sum scans of evens

totalEvens = broadcast last element

indx = constant array of 0..n

oddpos = totalEvens + indx - evpos

pos = if evens then evpos else oddpos

Scatter input using pos as index

Repeat with next bit to left until done

Application: Data Parallel Radix Sort

4	7	2	6	3	5	1	0
---	---	---	---	---	---	---	---

input

0	1	0	0	1	1	1	0
---	---	---	---	---	---	---	---

odds = last bit of each element

1	0	1	1	0	0	0	1
---	---	---	---	---	---	---	---

evens = complement of odds (last bit = 0)

0	1	1	2	3	3	3	4
---	---	---	---	---	---	---	---

evpos = exclusive sum scans of evens

totalEvens = broadcast last element

indx = constant array of 0..n

oddpos = totalEvens + indx - epos

pos = if evens then evpos else oddpos

Scatter input using pos as index

Repeat with next bit to left until done

Application: Data Parallel Radix Sort

4	7	2	6	3	5	1	0
---	---	---	---	---	---	---	---

input

0	1	0	0	1	1	1	0
---	---	---	---	---	---	---	---

odds = last bit of each element

1	0	1	1	0	0	0	1
---	---	---	---	---	---	---	---

evens = complement of odds (last bit = 0)

0	1	1	2	3	3	3	4
---	---	---	---	---	---	---	---

evpos = exclusive sum scans of evens

4	4	4	4	4	4	4	4
---	---	---	---	---	---	---	---

totalEvens = broadcast last element

indx = constant array of 0..n

oddpos = totalEvens + indx - evpos

pos = if evens then evpos else oddpos

Scatter input using pos as index

Repeat with next bit to left until done

Application: Data Parallel Radix Sort

4	7	2	6	3	5	1	0
---	---	---	---	---	---	---	---

input

0	1	0	0	1	1	1	0
---	---	---	---	---	---	---	---

odds = last bit of each element

1	0	1	1	0	0	0	1
---	---	---	---	---	---	---	---

evens = complement of odds (last bit = 0)

0	1	1	2	3	3	3	3	4
---	---	---	---	---	---	---	---	---

evpos = exclusive sum scans of evens

4	4	4	4	4	4	4	4
---	---	---	---	---	---	---	---

totalEvens = broadcast last element

0	1	2	3	4	5	6	7
---	---	---	---	---	---	---	---

indx = constant array of 0..n

oddpos = totalEvens + indx - evpos

pos = if evens then evpos else oddpos

Scatter input using pos as index

Repeat with next bit to left until done

Application: Data Parallel Radix Sort

4	7	2	6	3	5	1	0
---	---	---	---	---	---	---	---

input

0	1	0	0	1	1	1	0
---	---	---	---	---	---	---	---

odds = last bit of each element

1	0	1	1	0	0	0	1
---	---	---	---	---	---	---	---

evens = complement of odds (last bit = 0)

0	1	1	2	3	3	3	3	4
---	---	---	---	---	---	---	---	---

evpos = exclusive sum scans of evens

4	4	4	4	4	4	4	4
---	---	---	---	---	---	---	---

totalEvens = broadcast last element

0	1	2	3	4	5	6	7
---	---	---	---	---	---	---	---

indx = constant array of 0..n

4	4	5	5	5	6	7	8
---	---	---	---	---	---	---	---

4+0-0	4+1-1	4+2-1	4+3-2	4+4-3	4+5-3	4+6-3	4+7-3
-------	-------	-------	-------	-------	-------	-------	-------

oddpos = totalEvens + indx - evpos

pos = if evens then evpos else oddpos

Scatter input using pos as index

Repeat with next bit to left until done

Application: Data Parallel Radix Sort

4	7	2	6	3	5	1	0
---	---	---	---	---	---	---	---

input

0	1	0	0	1	1	1	0
---	---	---	---	---	---	---	---

odds = last bit of each element

1	0	1	1	0	0	0	1
---	---	---	---	---	---	---	---

evens = complement of odds (last bit = 0)

0	1	1	2	3	3	3	3	4
---	---	---	---	---	---	---	---	---

evpos = exclusive sum scans of evens

4	4	4	4	4	4	4	4
---	---	---	---	---	---	---	---

totalEvens = broadcast last element

0	1	2	3	4	5	6	7
---	---	---	---	---	---	---	---

indx = constant array of 0..n

4	4	5	5	5	6	7	8
---	---	---	---	---	---	---	---

4+0-0	4+1-1	4+2-1	4+3-2	4+4-3	4+5-3	4+6-3	4+7-3
-------	-------	-------	-------	-------	-------	-------	-------

oddpos = totalEvens + indx - evpos

0	4	1	2	5	6	7	3
---	---	---	---	---	---	---	---

pos = if evens then evpos else oddpos

Scatter input using pos as index

Repeat with next bit to left until done

Application: Data Parallel Radix Sort

4	7	2	6	3	5	1	0
---	---	---	---	---	---	---	---

input

0	1	0	0	1	1	1	0
---	---	---	---	---	---	---	---

odds = last bit of each element

1	0	1	1	0	0	0	1
---	---	---	---	---	---	---	---

evens = complement of odds (last bit = 0)

0	1	1	2	3	3	3	3	4
---	---	---	---	---	---	---	---	---

evpos = exclusive sum scans of evens

4	4	4	4	4	4	4	4
---	---	---	---	---	---	---	---

totalEvens = broadcast last element

0	1	2	3	4	5	6	7
---	---	---	---	---	---	---	---

indx = constant array of 0..n

4+0-0	4+1-1	4+2-1	4+3-2	4+4-3	4+5-3	4+6-3	4+7-3
-------	-------	-------	-------	-------	-------	-------	-------

oddpos = totalEvens + indx - evpos

*Using two
masked
assignments*

0	4	1	2	5	6	7	3
---	---	---	---	---	---	---	---

pos = if evens then evpos else oddpos

4	7	2	6	3	5	1	0
---	---	---	---	---	---	---	---

Scatter input using pos as index

Repeat with next bit to left until done

Application: Data Parallel Radix Sort

4	7	2	6	3	5	1	0
---	---	---	---	---	---	---	---

input

0	1	0	0	1	1	1	0
---	---	---	---	---	---	---	---

odds = last bit of each element

1	0	1	1	0	0	0	1
---	---	---	---	---	---	---	---

evens = complement of odds (last bit = 0)

0	1	1	2	3	3	3	4
---	---	---	---	---	---	---	---

evpos = exclusive sum scans of evens

4	4	4	4	4	4	4	4
---	---	---	---	---	---	---	---

totalEvens = broadcast last element

0	1	2	3	4	5	6	7
---	---	---	---	---	---	---	---

indx = constant array of 0..n

4+0-0	4+1-1	4+2-1	4+3-2	4+4-3	4+5-3	4+6-3	4+7-3
-------	-------	-------	-------	-------	-------	-------	-------

oddpos = totalEvens + indx - evpos

Using two
masked
assignments

0	4	1	2	5	6	7	3
---	---	---	---	---	---	---	---

pos = if evens then evpos else oddpos

4	7	2	6	3	5	1	0
---	---	---	---	---	---	---	---

Scatter input using pos as index

4	2	6	0	7	3	5	1
---	---	---	---	---	---	---	---

Repeat with next bit to left until done

Application: Data Parallel Radix Sort

4	7	2	6	3	5	1	0
---	---	---	---	---	---	---	---

input

0	1	0	0	1	1	1	0
---	---	---	---	---	---	---	---

odds = last bit of each element

1	0	1	1	0	0	0	1
---	---	---	---	---	---	---	---

evens = complement of odds (last bit = 0)

0	1	1	2	3	3	3	4
---	---	---	---	---	---	---	---

evpos = exclusive sum scans of evens

4	4	4	4	4	4	4	4
---	---	---	---	---	---	---	---

totalEvens = broadcast last element

0	1	2	3	4	5	6	7
---	---	---	---	---	---	---	---

indx = constant array of 0..n

4+0-0	4+1-1	4+2-1	4+3-2	4+4-3	4+5-3	4+6-3	4+7-3
-------	-------	-------	-------	-------	-------	-------	-------

Using two
masked
assignments

oddpos = totalEvens + indx - evpos

0	4	1	2	5	6	7	3
---	---	---	---	---	---	---	---

pos = if evens then evpos else oddpos

4	7	2	6	3	5	1	0
---	---	---	---	---	---	---	---

Scatter input using pos as index

4	2	6	0	7	3	5	1
---	---	---	---	---	---	---	---

Repeat with next bit to left until done

List Ranking with Pointer Doubling

Given a linked list of N nodes, find the distance (#hops) from each node to the end of the list.

Steps:

val = 1

while next != null

val += next.val

next =

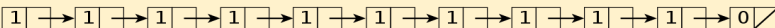
next.next

$d(n) =$

0 if $n.next$ is null

$1 + d(n.next)$ otherwise

Approach: put a processor at every node

Iteration 0 

Works if nodes
are on arbitrary
processors

List Ranking with Pointer Doubling

Given a linked list of N nodes, find the distance (#hops) from each node to the end of the list.

Steps:

val = 1

while next != null

val += next.val

next =

next.next


Works if nodes are on arbitrary processors

$d(n) =$

0 if $n.next$ is null

$1 + d(n.next)$ otherwise

Approach: put a processor at every node



List Ranking with Pointer Doubling

Given a linked list of N nodes, find the distance (#hops) from each node to the end of the list.

Steps:

val = 1

while next != null

val += next.val

next =

next.next


Works if nodes are on arbitrary processors

$d(n) =$

0 if $n.next$ is null

$1 + d(n.next)$ otherwise

Approach: put a processor at every node



List Ranking with Pointer Doubling

Given a linked list of N nodes, find the distance (#hops) from each node to the end of the list.

Steps:

val = 1

while next != null

val += next.val

next =

next.next


Works if nodes are on arbitrary processors

$d(n) =$

0 if $n.next$ is null

$1 + d(n.next)$ otherwise

Approach: put a processor at every node



List Ranking with Pointer Doubling

Given a linked list of N nodes, find the distance (#hops) from each node to the end of the list.

Steps:
val = 1
while next != null
 val += next.val
 next =
next.next


$d(n) =$

0 if $n.next$ is null

$1 + d(n.next)$ otherwise

Approach: put a processor at every node

Works if nodes
are on arbitrary
processors



Application: Fibonacci via Matrix Multiply Prefix

$$F_{n+1} = F_n + F_{n-1}$$

$$\begin{pmatrix} F_{n+1} \\ F_n \end{pmatrix} = \begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix} \begin{pmatrix} F_n \\ F_{n-1} \end{pmatrix}$$

Can compute all F_n by `matmul_prefix` on

$$\left[\begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix}, \begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix}, \begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix}, \begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix}, \begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix}, \begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix}, \begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix}, \begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix}, \begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix} \right]$$

then select the upper left entry

Application: Adding n -bit integers in $O(\log n)$ time

- Computing sum s of two n -bit binary numbers, think of a and b as array of bits
 - $a = a[n-1] a[n-2] \dots a[0]$ and $b = b[n-1] b[n-2] \dots b[0]$
 - $s = a+b = s[n] s[n-1] \dots s[0]$ (use carry-bit array $c = c[n-1] \dots c[0] c[-1]$)

Application: Adding n-bit integers in $O(\log n)$ time

- Computing sum s of two n -bit binary numbers, a and b
 - $a = a[n-1] a[n-2] \dots a[0]$ and $b = b[n-1] b[n-2] \dots b[0]$
 - $s = a+b = s[n] s[n-1] \dots s[0]$ (use carry-bit array $c = c[n-1] \dots c[0] c[-1]$)
- Formula
 - $c[-1] = 0$... rightmost carry bit
 - for $i = 0$ to $n-1$... compute right to left
 - $s[i] = (a[i] \text{ xor } b[i]) \text{ xor } c[i-1]$... one or three 1s
 - $c[i] = ((a[i] \text{ xor } b[i]) \text{ and } c[i-1]) \text{ or } (a[i] \text{ and } b[i])$... next carry bit

Application: Adding n-bit integers in $O(\log n)$ time

- Computing sum s of two n -bit binary numbers, a and b
 - $a = a[n-1] a[n-2] \dots a[0]$ and $b = b[n-1] b[n-2] \dots b[0]$
 - $s = a+b = s[n] s[n-1] \dots s[0]$ (use carry-bit array $c = c[n-1] \dots c[0] c[-1]$)
 - Formula
 - $c[-1] = 0$... rightmost carry bit
 - for $i = 0$ to $n-1$... compute right to left
 - $s[i] = (a[i] \text{ xor } b[i]) \text{ xor } c[i-1]$... one or three 1s
 - $c[i] = ((a[i] \text{ xor } b[i]) \text{ and } c[i-1]) \text{ or } (a[i] \text{ and } b[i])$... next carry bit
 - Example
 - $a = 22$
 - $b = 29$

$a =$	1	0	1	1	0		(22)
$b =$	1	1	1	0	1		(29)
$c =$	1	1	1	0	0	0	
$s =$	1	1	0	0	1	1	(51)
- Challenge: compute all $c[i]$ in $O(\log n)$ time via parallel prefix

Application: Adding n-bit integers in $O(\log n)$ time

- Recall carry bit calculation

$c[-1] = 0$... rightmost carry bit

for $i = 0$ to $n-1$

$c[i] = ((a[i] \text{ xor } b[i]) \text{ and } c[i-1]) \text{ or } (a[i] \text{ and } b[i])$... next carry bit

- Compute all $c[i]$ in $O(\log n)$ time via parallel prefix

Application: Adding n-bit integers in $O(\log n)$ time

- Recall carry bit calculation

$c[-1] = 0$... rightmost carry bit

for $i = 0$ to $n-1$

$c[i] = ((a[i] \text{ xor } b[i]) \text{ and } c[i-1]) \text{ or } (a[i] \text{ and } b[i])$... next carry bit

- Compute all $c[i]$ in $O(\log n)$ time via parallel prefix

for all $(0 \leq i \leq n-1)$ $p[i] = a[i] \text{ xor } b[i]$... propagate bit

for all $(0 \leq i \leq n-1)$ $g[i] = a[i] \text{ and } b[i]$... generate bit

Both $O(1)$ on
n procs

Application: Adding n-bit integers in $O(\log n)$ time

- Recall carry bit calculation

$c[-1] = 0$... rightmost carry bit

for $i = 0$ to $n-1$

$c[i] = ((a[i] \text{ xor } b[i]) \text{ and } c[i-1]) \text{ or } (a[i] \text{ and } b[i])$... next carry bit

- Compute all $c[i]$ in $O(\log n)$ time via parallel prefix

for all $(0 \leq i \leq n-1)$ $p[i] = a[i] \text{ xor } b[i]$... propagate bit

for all $(0 \leq i \leq n-1)$ $g[i] = a[i] \text{ and } b[i]$... generate bit

$$\begin{bmatrix} c[i] \\ 1 \end{bmatrix} = \begin{bmatrix} (p[i] \text{ and } c[i-1]) \text{ or } g[i] \\ 1 \end{bmatrix} = \begin{bmatrix} p[i] & g[i] \\ 0 & 1 \end{bmatrix} * \begin{bmatrix} c[i-1] \\ 1 \end{bmatrix} = M[i] * \begin{bmatrix} c[i-1] \\ 1 \end{bmatrix}$$

Application: Adding n-bit integers in $O(\log n)$ time

- Recall carry bit calculation

$c[-1] = 0$... rightmost carry bit

for $i = 0$ to $n-1$

$c[i] = ((a[i] \text{ xor } b[i]) \text{ and } c[i-1]) \text{ or } (a[i] \text{ and } b[i])$... next carry bit

- Compute all $c[i]$ in $O(\log n)$ time via parallel prefix

for all $(0 \leq i \leq n-1)$ $p[i] = a[i] \text{ xor } b[i]$... propagate bit

for all $(0 \leq i \leq n-1)$ $g[i] = a[i] \text{ and } b[i]$... generate bit

$$\begin{aligned} \begin{bmatrix} c[i] \\ 1 \end{bmatrix} &= \begin{bmatrix} (p[i] \text{ and } c[i-1]) \text{ or } g[i] \\ 1 \end{bmatrix} = \begin{bmatrix} p[i] & g[i] \\ 0 & 1 \end{bmatrix} * \begin{bmatrix} c[i-1] \\ 1 \end{bmatrix} = M[i] * \begin{bmatrix} c[i-1] \\ 1 \end{bmatrix} \\ &= M[i] * M[i-1] * \dots * M[0] * \begin{bmatrix} 0 \\ 1 \end{bmatrix} \end{aligned}$$

Application: Adding n-bit integers in $O(\log n)$ time

- Recall carry bit calculation

$c[-1] = 0$... rightmost carry bit

for $i = 0$ to $n-1$

$c[i] = ((a[i] \text{ xor } b[i]) \text{ and } c[i-1]) \text{ or } (a[i] \text{ and } b[i])$... next carry bit

- Compute all $c[i]$ in $O(\log n)$ time via parallel prefix

for all $(0 \leq i \leq n-1)$ $p[i] = a[i] \text{ xor } b[i]$... propagate bit

for all $(0 \leq i \leq n-1)$ $g[i] = a[i] \text{ and } b[i]$... generate bit

$$\begin{aligned} \begin{bmatrix} c[i] \\ 1 \end{bmatrix} &= \begin{bmatrix} (p[i] \text{ and } c[i-1]) \text{ or } g[i] \\ 1 \end{bmatrix} = \begin{bmatrix} p[i] & g[i] \\ 0 & 1 \end{bmatrix} * \begin{bmatrix} c[i-1] \\ 1 \end{bmatrix} = M[i] * \begin{bmatrix} c[i-1] \\ 1 \end{bmatrix} \\ &= M[i] * M[i-1] * \dots * M[0] * \begin{bmatrix} 0 \\ 1 \end{bmatrix} \end{aligned}$$

... evaluate $M[i] * M[i-1] * \dots * M[0]$ by parallel prefix

... 2-by-2 Boolean matrix multiplication is associative

Application: Adding n-bit integers in $O(\log n)$ time

- Recall carry bit calculation

$c[-1] = 0$... rightmost carry bit

for $i = 0$ to $n-1$

$c[i] = ((a[i] \text{ xor } b[i]) \text{ and } c[i-1]) \text{ or } (a[i] \text{ and } b[i])$... next carry bit

- Compute all $c[i]$ in $O(\log n)$ time via parallel prefix

for all $(0 \leq i \leq n-1)$ $p[i] = a[i] \text{ xor } b[i]$... propagate bit

for all $(0 \leq i \leq n-1)$ $g[i] = a[i] \text{ and } b[i]$... generate bit


$$\begin{aligned} \begin{bmatrix} c[i] \\ 1 \end{bmatrix} &= \begin{bmatrix} (p[i] \text{ and } c[i-1]) \text{ or } g[i] \\ 1 \end{bmatrix} = \begin{bmatrix} p[i] & g[i] \\ 0 & 1 \end{bmatrix} * \begin{bmatrix} c[i-1] \\ 1 \end{bmatrix} = M[i] * \begin{bmatrix} c[i-1] \\ 1 \end{bmatrix} \\ &= M[i] * M[i-1] * \dots * M[0] * \begin{bmatrix} 0 \\ 1 \end{bmatrix} \end{aligned}$$

... evaluate $M[i] * M[i-1] * \dots * M[0]$ by parallel prefix

... 2-by-2 Boolean matrix multiplication is associative

- Used in all computers to -- Carry look-ahead addition


This idea is used in all hardware
...Even going back to Babbage



Lexical analysis (tokenizing, scanning)

- Given a language of:
 - Identifiers (Z): string of chars
 - Strings (S): in double quotes
 - Ops (*): +, -, *, =, <, >, <=, >=
 - Expression (E), Quotes (Q),...
- Lexical analysis
 - Divide into tokens

Subset of Finite State Machine for Lexical Analysis



if	x	<=	n	then	print	"hello world"
----	---	----	---	------	-------	---------------

Full finite state machine encoded in table

Old State	Character Read													New line
•	A	B	...	Y	Z	+	-	*	<	>	=	"	Space	
N	A	A	...	A	A	*	*	*	<	<	*	Q	N	N
A	Z	Z	...	Z	Z	*	*	*	<	<	*	Q	N	N
Z	Z	Z	...	Z	Z	*	*	*	<	<	*	Q	N	N
*	A	A	...	A	A	*	*	*	<	<	*	Q	N	N
<	A	A	...	A	A	*	*	*	<	<	=	Q	N	N
=	A	A	...	A	A	*	*	*	<	<	*	Q	N	N
Q	S	S	...	S	S	S	S	S	S	S	S	E	S	S
S	S	S	...	S	S	S	S	S	S	S	S	E	S	S
E	E	E	...	E	E	*	*	*	<	<	*	S	N	N

- Each state in first column; **N** initial state.
- Each row gives the next state based on the next character at the top.
- Apply string **Y**+ to state **Z** written as **ZY**+ = ((**ZY**)")+ = (**Z**)"+ = **Q**+ = **S**
- Each column is a state transition for that character

Hillis and Steele, CACM 1986

Lexical analysis (tokenizing, scanning)

- Lexical analysis
 - Replace every character in the string with the array representation of its state-to-state function (column).
 - Perform a parallel-prefix operation with \oplus as the array composition. Each character becomes an array representing the state-to-state function for that prefix.
 - Use initial state (N, row 1) to index into these arrays.

	i	f		x	<	=		n	
N	A	A	N	A	N	<	*	N	A
A	Z	Z	N	Z	N	<	*	N	Z
Z	Z	Z	N	Z	N	<	*	N	Z
*	A	A	N	A	N	<	*	N	A
<	A	A	N	A	N	<	=	N	A
=	A	A	N	A	N	<	*	N	A
Q	S	S	S	S	S	S	S	S	S
S	S	S	S	S	S	S	S	S	S
E	E	E	N	E	N	<	*	N	E

<= n
A
A
A
A
A
A
S
S
A

Old State	Character Read													New line
	•	A	B	...	Y	Z	+	-	*	<	>	=	"	
N	A	A	...	A	A	*	*	*	<	<	*	Q	N	N
A	Z	Z	...	Z	Z	*	*	*	<	<	*	Q	N	N
Z	Z	Z	...	Z	Z	*	*	*	<	<	*	Q	N	N
*	A	A	...	A	A	*	*	*	<	<	*	Q	N	N
<	A	A	...	A	A	*	*	*	<	<	=	Q	N	N
=	A	A	...	A	A	*	*	*	<	<	*	Q	N	N
Q	S	S	...	S	S	S	S	S	S	S	S	E	S	S
S	S	S	...	S	S	S	S	S	S	S	S	E	S	S
E	E	E	...	E	E	*	*	*	<	<	*	S	N	N

Hillis and Steele, CACM 1986

Inverting triangular n-by-n matrices

in $O(\log^2 n)$ time

- Fact:
$$\begin{bmatrix} A & 0 \\ C & B \end{bmatrix}^{-1} = \begin{bmatrix} A^{-1} & 0 \\ -B^{-1}CA^{-1} & B^{-1} \end{bmatrix}$$

Inverting triangular n-by-n matrices

in $O(\log^2 n)$ time

- Fact: $\begin{bmatrix} A & 0 \\ C & B \end{bmatrix}^{-1} = \begin{bmatrix} A^{-1} & 0 \\ -B^{-1}CA^{-1} & B^{-1} \end{bmatrix}$
- Function `Tri_Inv(T)` // assume $n = \dim(T) = 2^m$ for simplicity

```
if T is 1-by-1  
    return 1/T
```

```
else  
    Write  $T = \begin{bmatrix} A & 0 \\ C & B \end{bmatrix}$ 
```

```
in parallel do {
```

```
    invA = Tri_Inv(A)
```

```
    invB = Tri_Inv(B)    // implicitly uses a tree
```

```
}
```

```
newC = -invB * C * invA //  $\log(n)$  for matmuls
```

```
return  $\begin{bmatrix} \text{invA} & 0 \\ \text{newC} & \text{invB} \end{bmatrix}$ 
```

Inverting triangular n-by-n matrices

in $O(\log^2 n)$ time

- Fact:
$$\begin{bmatrix} A & 0 \\ C & B \end{bmatrix}^{-1} = \begin{bmatrix} A^{-1} & 0 \\ -B^{-1}CA^{-1} & B^{-1} \end{bmatrix}$$

- Function `Tri_Inv(T)` // assume $n = \dim(T) = 2^m$ for simplicity

```
if T is 1-by-1
    return 1/T
```

```
else
```

Write $T = \begin{bmatrix} A & 0 \\ C & B \end{bmatrix}$

```
in parallel do {
```

```
    invA = Tri_Inv(A)
```

```
    invB = Tri_Inv(B)    // implicitly uses a tree
```

```
}
```

```
newC = -invB * C * invA // log(n) for matmuls
```

```
return  $\begin{bmatrix} \text{invA} & 0 \\ \text{newC} & \text{invB} \end{bmatrix}$ 
```

$$\begin{aligned} \text{time}(\text{Tri_Inv}(n)) &= \\ &\text{time}(\text{Tri_Inv}(n/2)) + O(\log(n)) \end{aligned}$$

Change variable to $m = \log n$ to
get $\text{time}(\text{Tri_Inv}(n)) = O(\log^2 n)$

Segmented Scans

Inputs = value array, flag array,
associative operator \oplus

Inclusive segmented sum scan

1	2	3	4	5	6	7	8
0	0	1	0	0	1	0	1

Flags are sometimes done with
Boolean and switch points

F	F	T	T	T	F	F	T
---	---	---	---	---	---	---	---


Result

1	3	3	7	12	6	13	8
---	---	---	---	----	---	----	---

Mapping Data Parallelism to Real Hardware

Connection Machine (CM-1,2)


Because communication is more important than processors



- Designed for AI by Thinking Machines Corporation (Hillis and Handler)
- CM-1 and CM-2 SIMD Design
 - 65,536 1-bit processors with 4 KB of memory each
 - 12-D boolean n-cube network (Feynman)
 - CM-2 add 1 floating point processor per 32 1-bit
- Programmed with data parallel languages
 - *Lisp
 - C*
- CM-5 was RISC+Vectors


SIMD/Vector Processors Use Data Parallelism

- SIMD instructions operate on a vector of elements
 - **These are specified as operations on vector registers**




- **Vectors “virtualize” the # of lanes (registers wider than #ALUs)**
- **SIMD on CPUs does not)**

(performs #pipes
adds in parallel)




SIMD/Vector Processors Use Data Parallelism

- SIMD instructions operate on a vector of elements
 - These are specified as operations on vector registers




- Vectors “virtualize” the # of lanes (registers wider than #ALUs)
- SIMD on CPUs does not)

(performs #pipes
adds in parallel)



SIMD/Vector Processors Use Data Parallelism

- SIMD instructions operate on a vector of elements
 - **These are specified as operations on vector registers**

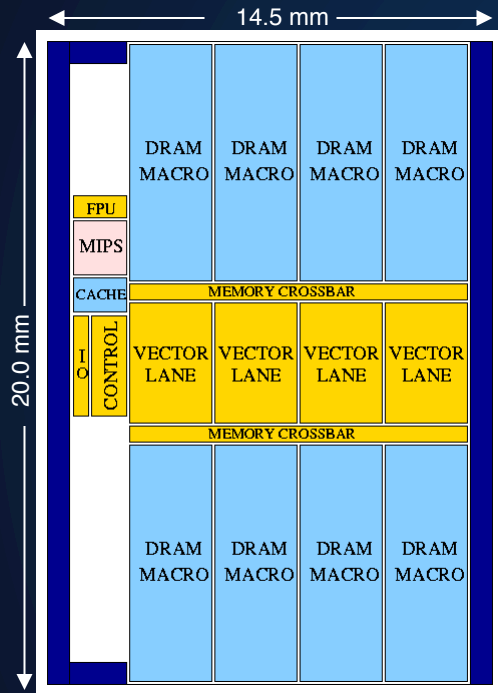


- **Vectors “virtualize” the # of lanes (registers wider than #ALUs)**
- **SIMD on CPUs does not**


(performs #pipes
adds in parallel)



VIRAM Processor at Berkeley


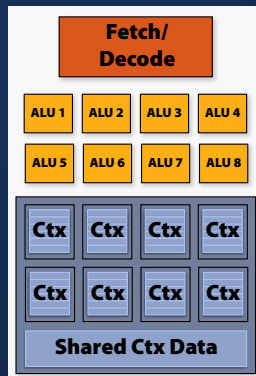
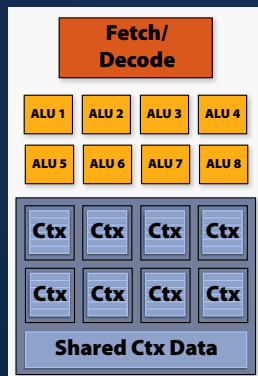



- MIPS Scalar core + 4-lane vector at 200 MHz (in 2002)
- 32-wide vector registers (in 64b)
- Peak vector performance
 - 1.6/3.2/6.4 Gops wo. multiply-add (64b/32b/16b operations)
 - 1.6 Gflops (single-precision)
- Transistor count: ~130M
- Reduction scan (in-register permutation instructions)



Mapping to GPUs

- For n-way parallelism may use n threads, divided into blocks
- Merge across statements (so $A=B$; $C=A$; is a single kernel)
- Mapping threads to ALUs and blocks to SMs is compiler / hardware problem



Bottom Line

- Branches are still expensive on GPUs
- May pad with zeros / nulls etc. to get length
- Often write code with a guard (if $i < n$), which will turn into mask – fine if n is large
- Non-contiguous memory is supported, but will still have a higher cost
- Enough parallelism to keep ALUs busy
and hide latency, memory/scheduling tradeoff

Mapping Data Parallelism to SMPs (and MPPs)

n-way parallelism onto *p*-way hardware

- Binary and unary operations

C = A+B



+



=



Cost $O(n/p)$

p speedup

- If arrays are not “aligned” then false sharing / communication require

Mapping Data Parallelism to SMPs (and MPPs)

n-way parallelism onto *p*-way hardware

- Binary and unary operations

C = A+B



+



=



Cost $O(n/p)$

p speedup


- If arrays are not “aligned” then false sharing / communication require
- Reductions and broadcasts



Cost $O(n/p)$

s = sum(C)

s:



tree reduction


+ log *p*)

Almost *p* speedup

Parallel prefix cost on p “big” processors

Time for this algorithm in parallel:

- $T_p(n) = O(n/p + \log p)$



Compute local
prefix sums in
 n/p steps

Updates across
processors in
 $\log p$ steps

serial time
on each
processor

communication and
computation up and
down the processor tree

The myth of $\log n$

- The $\log_2 n$ span is **not** the main reason for the usefulness of parallel prefix.

- Say $n = k \cdot p$ ($k = 1,000,000$ elements per proc)

$$\text{-- Cost} = \boxed{(k \text{ adds})} + \boxed{(\log_2 P \text{ steps})} + \boxed{(k \text{ adds})}$$

compute and store k
values $a[0]..a[k-1]$

parallel scan on
 $a[k-1]$ values

add 'my' scan result
to $a[0]..a[k-1]$

(2,000,000 local adds are serial for each processor, of course)

Key to implementing data parallel algorithms on clusters,
SMPs, MPPs, i.e., modern supercomputers

Summary of Data Parallelism

- Sequential semantics (or nearly) is very nice
 - Debugging is much easier without non-determinism
 - Correctness easier to reason about
- Cost model is independent of number of processors
 - How much inherent parallelism
- Need to “throttle” parallelism
 - $n \gg p$ can be hard to map, especially with nesting
 - Memory use is a problem
- More reading
 - Classic paper by Hillis and Steele “Data Parallel Algorithms”
<https://doi.org/10.1145/7902.7903> and on Youtube
 - Blelloch the NESL languages and “NESL Revisited paper, 2006