

---

# **CS 267**


## **Lecture 9: Distributed Memory Machines and Programming**

**Aydin Buluc**

**<https://sites.google.com/lbl.gov/cs267-spr2021/>**

- **Distributed Memory Architectures**
  - **Properties of communication networks**
  - **Topologies**
  - **Performance models**
- **Programming Distributed Memory Machines using Message Passing**
  - **Overview of MPI**
  - **Basic send/receive use**
  - **Non-blocking communication**
  - **Collectives**


# Architectures in Top 500 (systems share)



# Historical Perspective

---

- **Early distributed memory machines were:**
  - Collection of microprocessors.
  - Communication was performed using bi-directional queues between nearest neighbors.
- **Messages were forwarded by processors on path.**
  - “Store and forward” networking
- **There was a strong emphasis on topology in algorithms, in order to minimize the number of hops = minimize time**



## Network Analogy

---

**To have a large number of different transfers occurring at once, you need a large number of distinct wires**

- Not just a bus, as in shared memory

**Networks are like streets:**

- **Link** = street.
- **Switch** = intersection.
- **Distances** (hops) = number of blocks traveled.
- **Routing algorithm** = travel plan.

**Properties:**

**Latency:** how long to get between nodes in the network.

- Street: time for one car =  $\text{dist (miles)} / \text{speed (miles/hr)}$

**Bandwidth:** how much data can be moved per unit time.

- Street:  $\text{cars/hour} = \text{density (cars/mile)} * \text{speed (miles/hr)} * \text{\#lanes}$
- Network bandwidth is limited by the bit rate per wire and  $\text{\#wires}$

# Design Characteristics of a Network

---


- **Topology** (how things are connected)
  - Crossbar; ring; 2-D, 3-D, higher-D mesh or torus; hypercube; tree; butterfly; perfect shuffle, dragon fly, ...
- **Routing algorithm:**
  - Example in 2D torus: all east-west then all north-south (avoids deadlock).
- **Switching strategy:**
  - Circuit switching: full path reserved for entire message, like the telephone.
  - Packet switching: message broken into separately-routed packets, like the post office, or internet
- **Flow control** (what if there is congestion):
  - Stall, store data temporarily in buffers, re-route data to other nodes, tell source node to temporarily halt, discard, etc.

# Performance Properties of a Network: Latency

---

- **Diameter:** the maximum (over all pairs of nodes) of the shortest path between a given pair of nodes.
- **Latency:** delay between send and receive times
  - Latency tends to vary widely across architectures
  - Vendors often report **hardware latencies** (wire time)
  - Application programmers care about **software latencies** (user program to user program)
- **Observations:**
  - Latencies differ by 1-2 orders across network designs
  - Software/hardware overhead at source/destination dominate cost (1s-10s usecs)
  - Hardware latency varies with distance (10s-100s nsec per hop) but is small compared to overheads
- **Latency is key for programs with many small messages**

# End to End Latency (1/2 roundtrip) Over Time



- Latency has not improved significantly, unlike Moore's Law




# Performance Properties of a Network: Bandwidth


---

- The **bandwidth** of a link = # wires / time-per-bit
- Bandwidth typically in Gigabytes/sec (GB/s), i.e.,  $8 \times 2^{20}$  bits per second
- **Effective bandwidth** is usually lower than physical link bandwidth due to packet overhead.

- Bandwidth is important for applications with mostly large messages




# Bandwidth Chart



Data from Mike Welcome, NERSC

# Performance Properties of a Network: Bisection Bandwidth

- **Bisection bandwidth:** bandwidth across smallest cut that divides network into two equal halves
- Bandwidth across “narrowest” part of the network



- Bisection bandwidth is important for algorithms in which all processors need to communicate with all others

# Linear and Ring Topologies



---

## ◦ Linear array



- Diameter =  $n-1$ ; average distance  $\sim n/3$ .
- Bisection bandwidth = 1 (in units of link bandwidth).

## ◦ Torus or Ring




- Diameter =  $n/2$ ; average distance  $\sim n/4$ .
- Bisection bandwidth = 2.
- Natural for algorithms that work with 1D arrays.

# Meshes and Tori – used in Hopper

---


## Two dimensional mesh

- Diameter =  $2 * (\text{sqrt}(n) - 1)$
- Bisection bandwidth =  $\text{sqrt}(n)$



## Two dimensional torus

- Diameter =  $\text{sqrt}(n)$
- Bisection bandwidth =  $2 * \text{sqrt}(n)$




- Generalizes to higher dimensions
  - Cray XT (eg Hopper@NERSC) uses 3D Torus
- Natural for algorithms that work with 2D and/or 3D arrays (matmul)

# Hypercubes

◦ Number of nodes  $n = 2^d$  for dimension  $d$ .

- Diameter =  $d$ .
- Bisection bandwidth =  $n/2$ .




◦ Popular in early machines (Intel iPSC, NCUBE).

- Lots of clever algorithms.
- See 1996 online CS267 notes.




◦ Greycode addressing:

- Each node connected to  $d$  others with 1 bit different.




# Trees

- **Diameter =  $\log n$ .**
- **Bisection bandwidth = 1.**
- **Easy layout as planar graph.**
- **Many tree algorithms (e.g., summation).**
- **Fat trees avoid bisection bandwidth problem:**
  - More (or wider) links near top.
  - Example: Thinking Machines CM-5.





# Butterflies

- Really an unfolded version of hypercube.
- A d-dimensional butterfly has  $(d+1) 2^d$  "**switching nodes**" (not to be confused with processors, which is  $n = 2^d$ )
- Butterfly was invented because hypercube required increasing radix of switches as the network got larger; prohibitive at the time
- Diameter =  $\log n$ . Bisection bandwidth =  $n$
- No path diversity: bad with adversarial traffic



A row of butterfly is a node in hypercube




**butterfly switch**

**Ex: to get from proc 101 to 110,  
Compare bit-by-bit and  
Switch if they disagree, else not**



# Does Topology Matter?



See EECS Tech Report *UCB/EECS-2011-92*, August 2011

# Why so many topologies?

---

- **Different systems have different needs**
  - Size of the system (data center vs. NIC)
- **Complexity vs. optimality**
- **Physical constraints**
  - Innovations in HW enable previously infeasible technologies
- **Two recent technological changes:**
  - Higher radix (number of ports supported) switches economical, which is really a consequence of Moore's law
  - Fiber optic is feasible → distance doesn't matter

# Dragonflies – used in Edison and Cori


---

- **Motivation:** Exploit gap in cost and performance between optical interconnects (which go between cabinets in a machine room) and electrical networks (inside cabinet)
  - Optical (fiber) more expensive but higher bandwidth when long
  - Electrical (copper) networks cheaper, faster when short
- **Combine in hierarchy:**
  - Several groups are connected together using all to all links, i.e. each group has at least one link directly to each other group.
  - The topology inside each group can be any topology.
- **Uses a randomized routing algorithm**
- **Outcome:** programmer can (usually) ignore topology, get good performance
  - Important in virtualized, dynamic environment
  - Drawback: variable performance

“Technology-Drive, Highly-Scalable Dragonfly Topology,” ISCA 2008


# Dragonfly in practice

Aries interconnect for an 8 cabinet Cray XC40

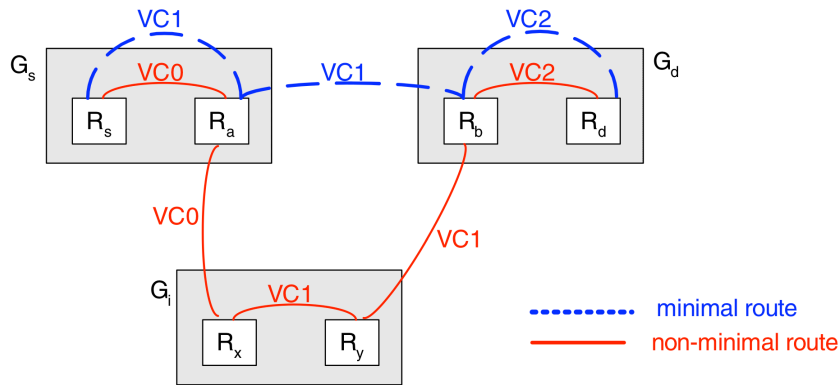


Source: European Centre for Medium-Range Weather Forecasts

Source of image on the right (and more info):  
<http://www.nersc.gov/users/computational-systems/edison/configuration/interconnect/>



# Why randomized routing?



## Minimal routing:

1. If  $G_s \neq G_d$  and  $R_s$  has no connection to  $G_d$ , route within  $G_s$  from  $R_s$  to  $R_a$ , which has a global channel to  $G_d$ .
2. If  $G_s \neq G_d$ , traverse the global channel from  $R_a$  to router  $R_b$  in  $G_d$ .
3. If  $R_b \neq R_d$ , route from  $R_b$  to  $R_d$ .

Minimal routing works well when things are load balanced, potentially catastrophic in adversarial traffic patterns.

**Randomization idea:** For each packet sourced at router  $R_s \in G_s$  and addressed to a router in another group  $R_d \in G_d$ , first route it to an intermediate group  $G_i$ .

- This requires at most two group-level link traversals
- And at most 5 total link traversals

Valiant, Leslie G. "A scheme for fast parallel communication." *SIAM journal on computing* 11.2 (1982): 350-361.

---

# Performance Models

# Shared Memory Performance Models

---

- **Parallel Random Access Memory (PRAM)**
- **All memory access operations complete in one clock period -- no concept of memory hierarchy (“too good to be true”).**
  - OK for understanding whether an algorithm has enough parallelism at all (see CS273).
  - Parallel algorithm design strategy: first do a PRAM algorithm, then worry about memory/communication time (sometimes works)
- **Slightly more realistic versions exist**
  - E.g., Concurrent Read Exclusive Write (CREW) PRAM.
  - Still missing the memory hierarchy

## Latency and Bandwidth Model

---

- Time to send message of length  $n$  is roughly

$$\begin{aligned}\text{Time} &= \text{latency} + n * \text{cost\_per\_word} \\ &= \text{latency} + n / \text{bandwidth}\end{aligned}$$

- Topology is assumed irrelevant.
- Often called “ $\alpha$ – $\beta$  model” and written

$$\text{Time} = \alpha + n * \beta$$

- Usually  $\alpha \gg \beta \gg \text{time per flop}$ .
  - One long message is cheaper than many short ones.
  - Can do hundreds or thousands of flops for cost of one message.


$$\alpha + n * \beta \ll n * (\alpha + 1 * \beta)$$

- Lesson: Need large computation-to-communication ratio to be efficient.
- LogP – more detailed model (Latency/overhead/gap/Proc.)




# Latency and Bandwidth Model in 2010 and beyond+


**Processors are multi-core and many nodes are multi-chip**



# NIC bandwidth bottleneck



- Ping-pong between 2 nodes using 1-16 cores on each node
- Top: BG/Q,
- Bottom Cray XE6
- X-axis is message size



- Alpha-beta model predicts a single curve – rates independent of the number of communicating processes

---

# Programming Distributed Memory Machines with Message Passing

Slides by

Aydin Buluc, Jonathan Carter, Jim Demmel,  
Bill Gropp, Kathy Yelick

# Message Passing Libraries (1)

---

- Many “message passing libraries” were once available
  - Chameleon, from ANL.
  - CMMD, from Thinking Machines.
  - Express, commercial.
  - MPL, native library on IBM SP-2.
  - NX, native library on Intel Paragon.
  - Zipcode, from LLL.
  - PVM, Parallel Virtual Machine, public, from ORNL/UTK.
  - Others...
  - **MPI, Message Passing Interface, now the industry standard.**
- Need standards to write portable code.

# Message Passing Libraries (2)

---

- **All communication, synchronization require subroutine calls**
  - **No shared variables**
  - **Program run on a single processor just like any uniprocessor program, except for calls to message passing library**
- **Subroutines for**
  - **Communication**
    - **Pairwise or point-to-point: Send and Receive**
    - **Collectives all processor get together to**
      - Move data: Broadcast, Scatter/gather
      - Compute and move: sum, product, max, prefix sum, ... of data on many processors
  - **Synchronization**
    - **Barrier**
    - **No locks because there are no shared variables to protect**
  - **Enquiries**
    - **How many processes? Which one am I? Any messages waiting?**

## Novel Features of MPI

---

- Communicators encapsulate communication spaces for library safety
- Datatypes reduce copying costs and permit heterogeneity
- Multiple communication modes allow precise buffer management
- Extensive collective operations for scalable global communication
- Process topologies permit efficient process placement, user views of process layout
- Profiling interface encourages portable tools

## MPI References




---

- **The Standard itself:**
  - at <http://www.mpi-forum.org>
  - All MPI official releases, in both postscript and HTML
  - Latest version MPI 3.1, released June 2015
- **Other information on Web:**
  - at <http://www.mcs.anl.gov/research/projects/mpi/index.htm>
  - pointers to lots of stuff, including other talks and tutorials, a FAQ, other MPI pages

# Books on MPI

---

- ***Using MPI: Portable Parallel Programming with the Message-Passing Interface (third edition)***, by Gropp, Lusk, and Skjellum, MIT Press, 2014.
- ***Using Advanced MPI: Modern Features of the Message-Passing Interface***, by Gropp, Hoefler, Thakur, and Lusk, MIT Press, 2014
- ***Using MPI-2: Portable Parallel Programming with the Message-Passing Interface***, by Gropp, Lusk, and Thakur, MIT Press, 1999.
- ***MPI: The Complete Reference - Vol 1 The MPI Core***, by Snir, Otto, Huss-Lederman, Walker, and Dongarra, MIT Press, 1998.
- ***MPI: The Complete Reference - Vol 2 The MPI Extensions***, by Gropp, Huss-Lederman, Lumsdaine, Lusk, Nitzberg, Saphir, and Snir, MIT Press, 1998.
- ***Designing and Building Parallel Programs***, by Ian Foster, Addison-Wesley, 1995.
- ***Parallel Programming with MPI***, by Peter Pacheco, Morgan-Kaufmann, 1997.





## Finding Out About the Environment

---

- Two important questions that arise early in a parallel program are:
  - How many processes are participating in this computation?
  - Which one am I?
- MPI provides functions to answer these questions:
  - `MPI_Comm_size` reports the number of processes.
  - `MPI_Comm_rank` reports the *rank*, a number between 0 and size-1, identifying the calling process

# Hello (C)

---

```
#include "mpi.h"
#include <stdio.h>

int main( int argc, char *argv[] )
{
    int rank, size;
    MPI_Init( &argc, &argv );
    MPI_Comm_rank( MPI_COMM_WORLD, &rank );
    MPI_Comm_size( MPI_COMM_WORLD, &size );
    printf( "I am %d of %d\n", rank, size );
    MPI_Finalize();
    return 0;
}
```

## Notes on Hello World


---

- All MPI programs begin with `MPI_Init` and end with `MPI_Finalize`
- `MPI_COMM_WORLD` is defined by `mpi.h` (in C) or `mpif.h` (in Fortran) and designates all processes in the MPI “job”
- Each statement executes independently in each process
  - including the `printf/print` statements
- The MPI-1 Standard does not specify how to run an MPI program, but many implementations provide `mpirun -np 4 a.out`

# MPI Basic Send/Receive

---

- **We need to fill in the details in**



- **Things that need specifying:**
  - How will “data” be described?
  - How will processes be identified?
  - How will the receiver recognize/screen messages?
  - What will it mean for these operations to complete?

## Some Basic Concepts

---

- Processes can be collected into groups
- Each message is sent in a context, and must be received in the same context
  - Provides necessary support for libraries
- A group and context together form a communicator
- A process is identified by its rank in the group associated with a communicator
- There is a default communicator whose group contains all initial processes, called `MPI_COMM_WORLD`

# MPI Datatypes

---

- The data in a message to send or receive is described by a triple (address, count, datatype), where
- An MPI datatype is recursively defined as:
  - predefined, corresponding to a data type from the language (e.g., MPI\_INT, MPI\_DOUBLE)
  - a contiguous array of MPI datatypes
  - a strided block of datatypes
  - an indexed array of blocks of datatypes
  - an arbitrary structure of datatypes
- There are MPI functions to construct custom datatypes, in particular ones for subarrays
- May hurt performance if datatypes are complex


## MPI Tags

---

- Messages are sent with an accompanying user-defined integer tag, to assist the receiving process in identifying the message
- Messages can be screened at the receiving end by specifying a specific tag, or not screened by specifying `MPI_ANY_TAG` as the tag in a receive
- Some non-MPI message-passing systems have called tags “message types”. MPI calls them tags to avoid confusion with datatypes

## MPI Basic (Blocking) Send

---



`MPI_Send( A, 10, MPI_DOUBLE, 1, ...)`

`MPI_Recv( B, 20, MPI_DOUBLE, 0, ... )`


**`MPI_SEND(start, count, datatype, dest, tag, comm)`**

- The message buffer is described by (`start`, `count`, `datatype`).
- The target process is specified by `dest`, which is the rank of the target process in the communicator specified by `comm`.
- When this function returns, the data has been delivered to the system and the buffer can be reused. The message may not have been received by the target process.



# MPI Basic (Blocking) Receive

---



`MPI_Send( A, 10, MPI_DOUBLE, 1, ...)`

`MPI_Recv( B, 20, MPI_DOUBLE, 0, ... )`

**`MPI_RECV(start, count, datatype, source, tag, comm, status)`**

- Waits until a matching (both `source` and `tag`) message is received from the system, and the buffer can be used
- `source` is rank in communicator specified by `comm`, or `MPI_ANY_SOURCE`
- `tag` is a tag to be matched or `MPI_ANY_TAG`
- receiving fewer than `count` occurrences of `datatype` is OK, but receiving more is an error
- `status` contains further information (e.g. size of message)

# A Simple MPI Program

---

```
#include "mpi.h"
#include <stdio.h>
int main( int argc, char *argv[])
{
    int rank, buf;
    MPI_Status status;
    MPI_Init(&argv, &argc);
    MPI_Comm_rank( MPI_COMM_WORLD, &rank );

    /* Process 0 sends and Process 1 receives */
    if (rank == 0) {
        buf = 123456;
        MPI_Send( &buf, 1, MPI_INT, 1, 0, MPI_COMM_WORLD);
    }
    else if (rank == 1) {
        MPI_Recv( &buf, 1, MPI_INT, 0, 0, MPI_COMM_WORLD,
                  &status );
        printf( "Received %d\n", buf );
    }

    MPI_Finalize();
    return 0;
}
```

## Retrieving Further Information

---

- **Status** is a data structure allocated in the user's program.
- **In C:**

```
int recvd_tag, recvd_from, recvd_count;
MPI_Status status;
MPI_Recv(..., MPI_ANY_SOURCE, MPI_ANY_TAG, ..., &status )
recvd_tag  = status.MPI_TAG;
recvd_from = status.MPI_SOURCE;
MPI_Get_count( &status, datatype, &recvd_count );
```

## MPI can be simple

---

- **Claim: most MPI applications can be written with only 6 functions (although which 6 may differ)**
- Using point-to-point:
  - `MPI_INIT`
  - `MPI_FINALIZE`
  - `MPI_COMM_SIZE`
  - `MPI_COMM_RANK`
  - `MPI_SEND`
  - `MPI_RECEIVE`
- Using collectives:
  - `MPI_INIT`
  - `MPI_FINALIZE`
  - `MPI_COMM_SIZE`
  - `MPI_COMM_RANK`
  - `MPI_BCAST`
  - `MPI_REDUCE`
- **You may use more for convenience or performance**

# But is that small subset the practical usage?

- 35 What aspects of the MPI standard do you use in your application in its current form?  
\* (multiple)
- 36 What aspects of the MPI standard appear in performance-critical sections of your current application? \* (multiple)
- 37 What aspects of the MPI standard do you anticipate using in the "exascale" version of your application? \* (multiple)


Responses	Q35: Current Usage			Q37: Exascale Usage			Q36: Performance Critical		
	AD	ST	Overall	AD	ST	Overall	AD	ST	Overall
Point-to-point communications	96%	79%	88%	89%	71%	80%	93%	75%	84%
MPI derived datatypes	25%	21%	23%	21%	21%	21%	14%	7%	11%
Collective communications	86%	75%	80%	96%	68%	82%	64%	64%	64%
Neighbor collective communications	14%	14%	14%	32%	25%	29%	7%	11%	9%
Communicators and group management	68%	54%	61%	61%	50%	55%	29%	7%	18%
Process topologies	14%	7%	11%	32%	11%	21%	4%	4%	4%
RMA (one-sided communications)	36%	7%	21%	50%	36%	43%	21%	7%	14%
RMA shared windows	18%	7%	12%	21%	18%	20%	7%	7%	7%
MPI I/O (called directly)	25%	18%	21%	21%	18%	20%	4%	7%	5%
MPI I/O (called through a third-party library)	32%	21%	27%	36%	25%	30%	7%	11%	9%
MPI profiling interface	11%	0%	14%	11%	21%	16%	0%	4%	2%

## But is that small subset the practical usage?

No.	Question and Responses	AD	ST	Overall
38	<b>What is the dominant communication in your application?</b> Check all that apply, recognizing that many applications have different communication patterns in different phases. ( <i>multiple</i> )			
	Each process talks to (almost) every other process	25%	21%	23%
	Processes communicate in fixed "neighborhoods" of limited size	46%	46%	46%
	Processes communicate in "neighborhoods" of limited size that may change in different phases of the application or evolve over the course of a run	42%	36%	39%
	Communication is largely irregular	18%	21%	20%
39	<b>Can your application take advantage of non-blocking point-to-point operations...</b> ( <i>multiple</i> )			
	To overlap communication with computation?	89%	71%	80%
	To allow asynchronous progress?	64%	64%	64%
	To allow event-based programming?	43%	29%	36%
40	<b>Can your application take advantage of non-blocking collective operations...</b> ( <i>multiple</i> )			
	To overlap communication with computation?	71%	46%	59%
	To allow asynchronous progress?	46%	29%	38%
	To allow event-based programming?	25%	14%	20%

<https://info.ornl.gov/sites/publications/Files/Pub108588.pdf>

# PI redux: Numerical integration



Mathematically, we know that:

$$\int_0^1 \frac{4.0}{(1+x^2)} dx = \pi$$

We can approximate the integral as a sum of rectangles:


$$\sum_{i=0}^N F(x_i) \Delta x \approx \pi$$

Where each rectangle has width  $\Delta x$  and height  $F(x_i)$  at the middle of interval  $i$ .

## Example: Calculating Pi

---

**E.g., in a 4-process run, each process gets every 4<sup>th</sup> interval. Process 0 slices are in red.**



- **Simple program written in a data parallel style in MPI**
  - E.g., for a reduction (recall “data parallelism” lecture), each process will first reduce (sum) its own values, then call a collective to combine them
- **Estimates pi by approximating the area of the quadrant of a unit circle**
- **Each process gets  $1/p$  of the intervals (mapped round robin, i.e., a cyclic mapping)**



## Example: PI in C – 1/2

---

```
#include "mpi.h"
#include <math.h>
#include <stdio.h>

int main(int argc, char *argv[])
{
    int done = 0, n, myid, numprocs, i, rc;
    double PI25DT = 3.141592653589793238462643;
    double mypi, pi, h, sum, x, a;
    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &numprocs);
    MPI_Comm_rank(MPI_COMM_WORLD, &myid);
    while (!done) {
        if (myid == 0) {
            printf("Enter the number of intervals: (0 quits) ");
            scanf("%d", &n);
        }
        MPI_Bcast(&n, 1, MPI_INT, 0, MPI_COMM_WORLD);
        if (n == 0) break;
    }
}
```

## Example: PI in C – 2/2

---

```
h    = 1.0 / (double) n;
sum = 0.0;
for (i = myid + 1; i <= n; i += numprocs) {
    x = h * ((double)i - 0.5);
    sum += 4.0 * sqrt(1.0 - x*x);
}
mypi = h * sum;
MPI_Reduce(&mypi, &pi, 1, MPI_DOUBLE, MPI_SUM, 0,
           MPI_COMM_WORLD);
if (myid == 0)
    printf("pi is approximately %.16f, Error is .16f\n",
           pi, fabs(pi - PI25DT));
}
MPI_Finalize();


return 0;

}
```

# Buffers

---


- When you send data, where does it go? One possibility is:



## Avoiding Buffering

---

- Avoiding copies uses less memory
- May use more or less time



This requires that `MPI_Send` wait on delivery, or that `MPI_Send` return before transfer is complete, and we wait later.

# Blocking and Non-blocking Communication

---

- So far we have been using *blocking* communication:
  - `MPI_Recv` does not complete until the buffer is full (available for use).
  - `MPI_Send` does not complete until the buffer is empty (available for use).
- Completion depends on size of message and amount of system buffering.

# Sources of Deadlocks

---

- **Send a large message from process 0 to process 1**
  - If there is insufficient storage at the destination, the send must wait for the user to provide the memory space (through a receive)
- **What happens with this code?**

Process 0

Process 1

---

**Send (1)**

**Send (0)**

**Recv (1)**

**Recv (0)**

- This is called “unsafe” because it depends on the availability of system buffers in which to store the data sent until it can be received

## Some Solutions to the “unsafe” Problem

---

- **Order the operations more carefully:**

Process 0

Process 1

---

**Send (1)**

**Recv (0)**

**Recv (1)**

**Send (0)**

- Supply receive buffer at same time as send:

Process 0

Process 1

---

**Sendrecv (1)**

**Sendrecv (0)**

## More Solutions to the “unsafe” Problem

---

- **Supply own space as buffer for send**

Process 0

Process 1

---

**Bsend(1)**

**Bsend(0)**

**Recv(1)**

**Recv(0)**

- **Use non-blocking operations:**

Process 0

Process 1

---

**Isend(1)**

**Isend(0)**

**Irecv(1)**

**Irecv(0)**

**Waitall**

**Waitall**



## MPI's Non-blocking Operations

---

- Non-blocking operations return (immediately) “request handles” that can be tested and waited on:

```
MPI_Request request;
```

```
MPI_Status status;
```

```
MPI_Isend(start, count, datatype,  
          dest, tag, comm, &request);
```

```
MPI_Irecv(start, count, datatype,  
          dest, tag, comm, &request);
```

```
MPI_Wait(&request, &status);
```

(each request must be Waited on)

- One can also test without waiting:

```
MPI_Test(&request, &flag, &status);
```

- Accessing the data buffer without waiting is undefined

## Multiple Completions

---

- It is sometimes desirable to wait on multiple requests:

```
MPI_Waitall(count, array_of_requests,  
            array_of_statuses)
```

```
MPI_Waitany(count, array_of_requests,  
            &index, &status)
```

```
MPI_Waitsome(count, array_of_requests,  
             array_of_indices, array_of_statuses)
```

- There are corresponding versions of test for each of these.

# Communication Modes

---

- **MPI provides multiple *modes* for sending messages:**
  - **Synchronous mode (MPI\_Ssend):** the send does not complete until a matching receive has begun. (Unsafe programs deadlock.)
  - **Buffered mode (MPI\_Bsend):** the user supplies a buffer to the system for its use. (User allocates enough memory to make an unsafe program safe.)
  - **Ready mode (MPI\_Rsend):** user guarantees that a matching receive has been posted.
    - Allows access to fast protocols
    - undefined behavior if matching receive not posted
- **Non-blocking versions (MPI\_Issend, etc.)**
- **MPI\_Recv receives messages sent in any mode.**
- **See [www.mpi-forum.org](http://www.mpi-forum.org) for summary of all flavors of send/receive**